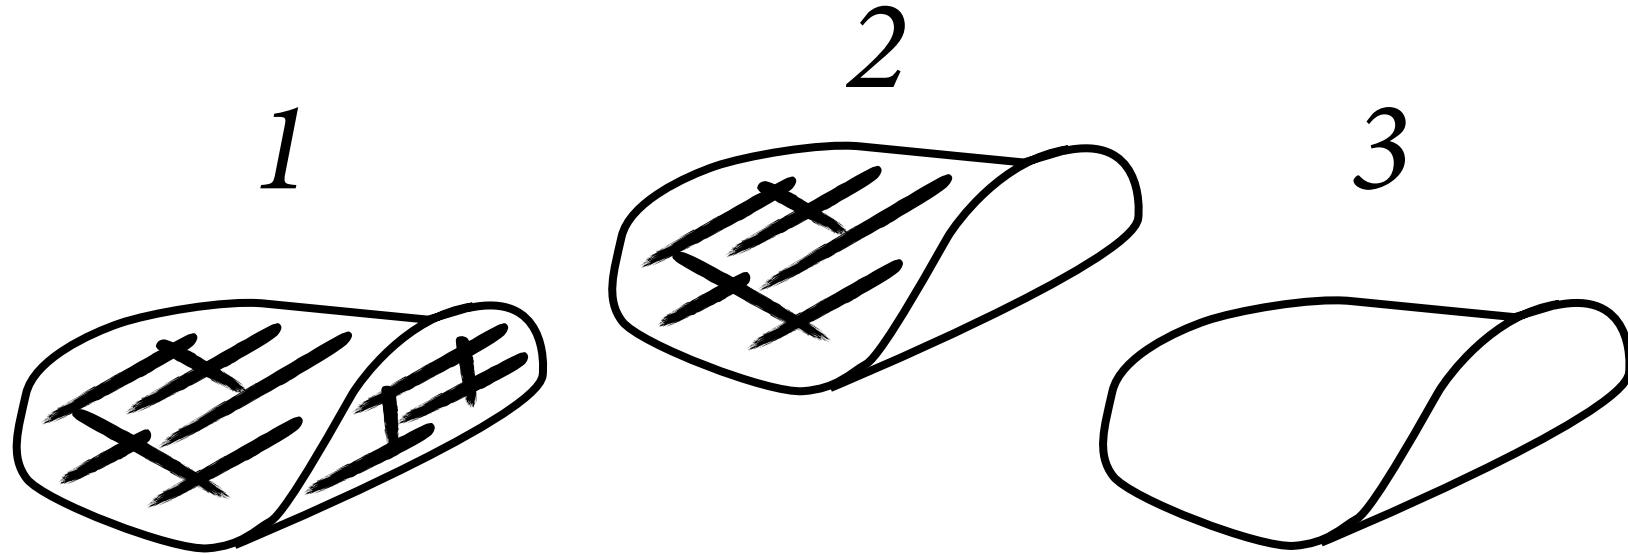


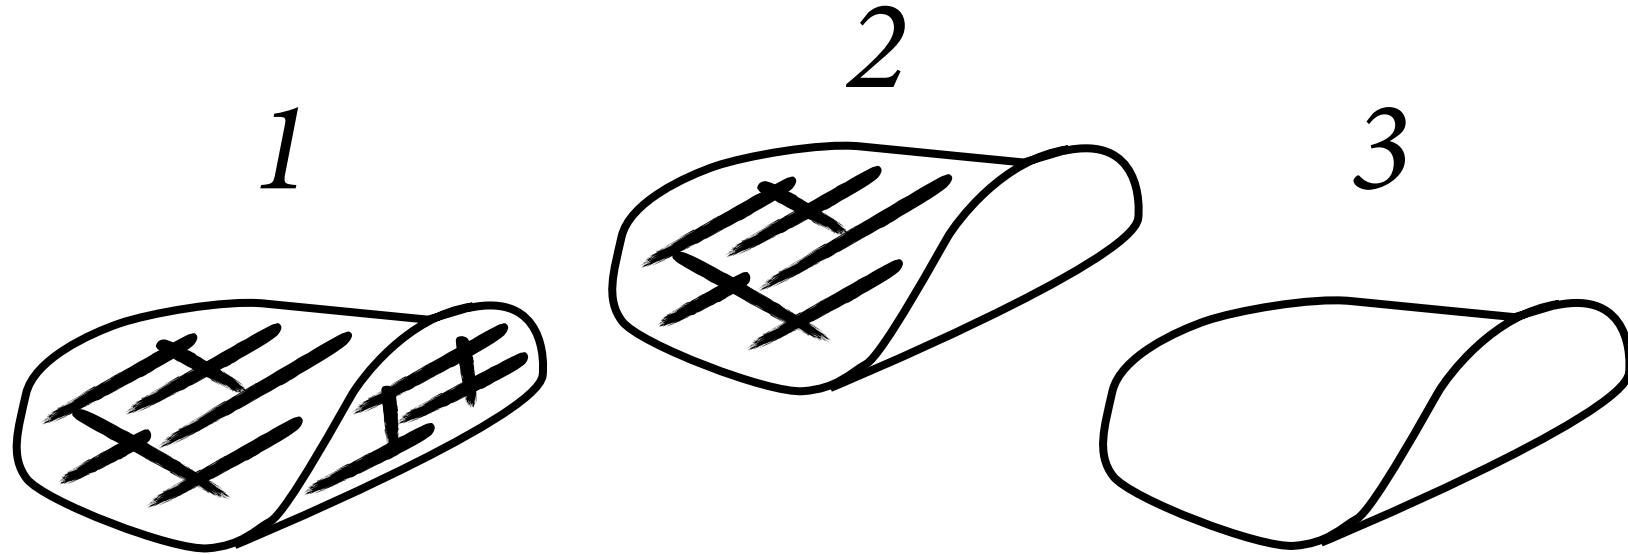
Statistical Rethinking

Winter 2019

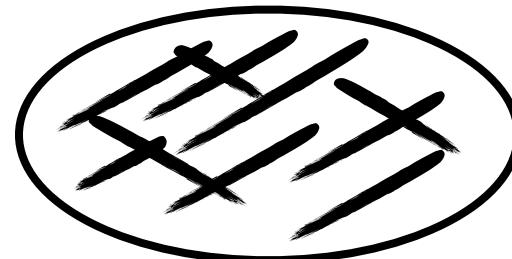
Lecture 20 / Week 10

Missing Data &
Other Opportunities





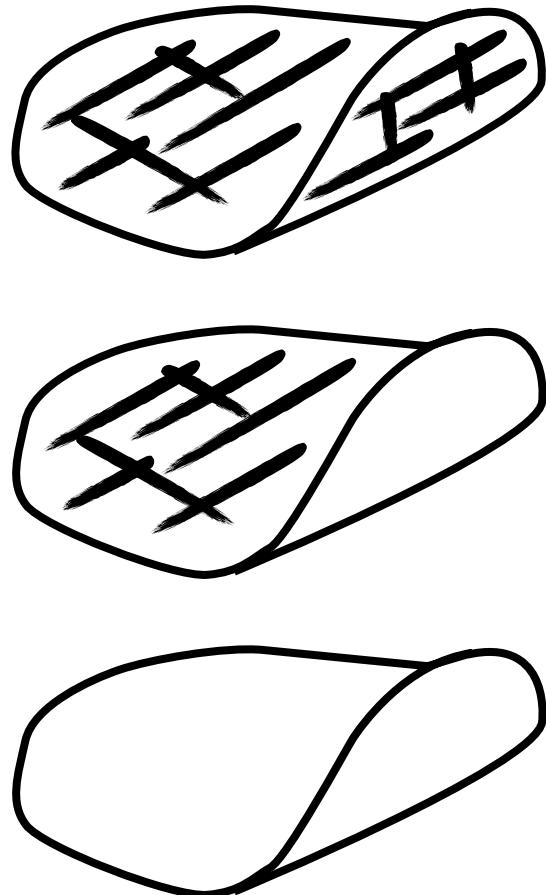
You are served:

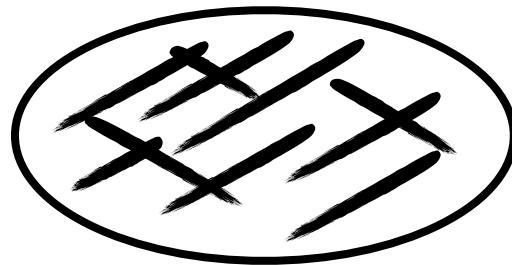


Probability other side is burnt?

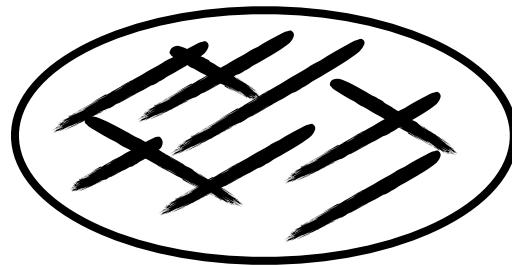
Avoid being clever

- Intuition terrible guide to probability
- No need to be clever; just ruthlessly apply conditional probability
 - $\Pr(\text{want to know} \mid \text{already know})$



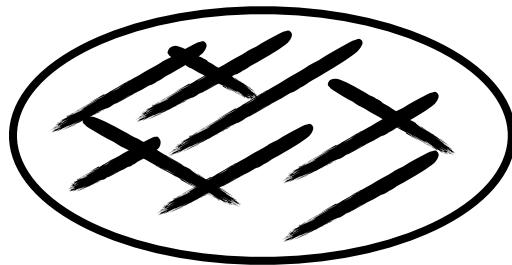


$$\Pr(\text{burnt down} | \text{burnt up}) = \frac{\Pr(\text{burnt up, burnt down})}{\Pr(\text{burnt up})}$$



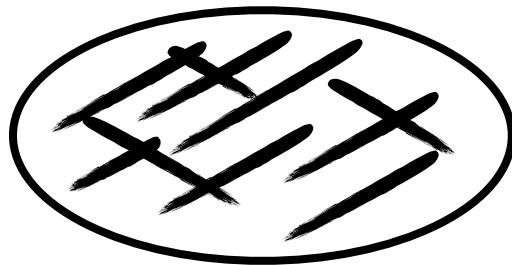
$$\Pr(\text{burnt down}|\text{burnt up}) = \frac{\Pr(\text{burnt up, burnt down})}{\Pr(\text{burnt up})}$$

$$\Pr(\text{burnt up}) = \Pr(\text{BB})(1) + \Pr(\text{BU})(0.5) + \Pr(\text{UU})(0)$$



$$\Pr(\text{burnt down} | \text{burnt up}) = \frac{\Pr(\text{burnt up, burnt down})}{\Pr(\text{burnt up})}$$

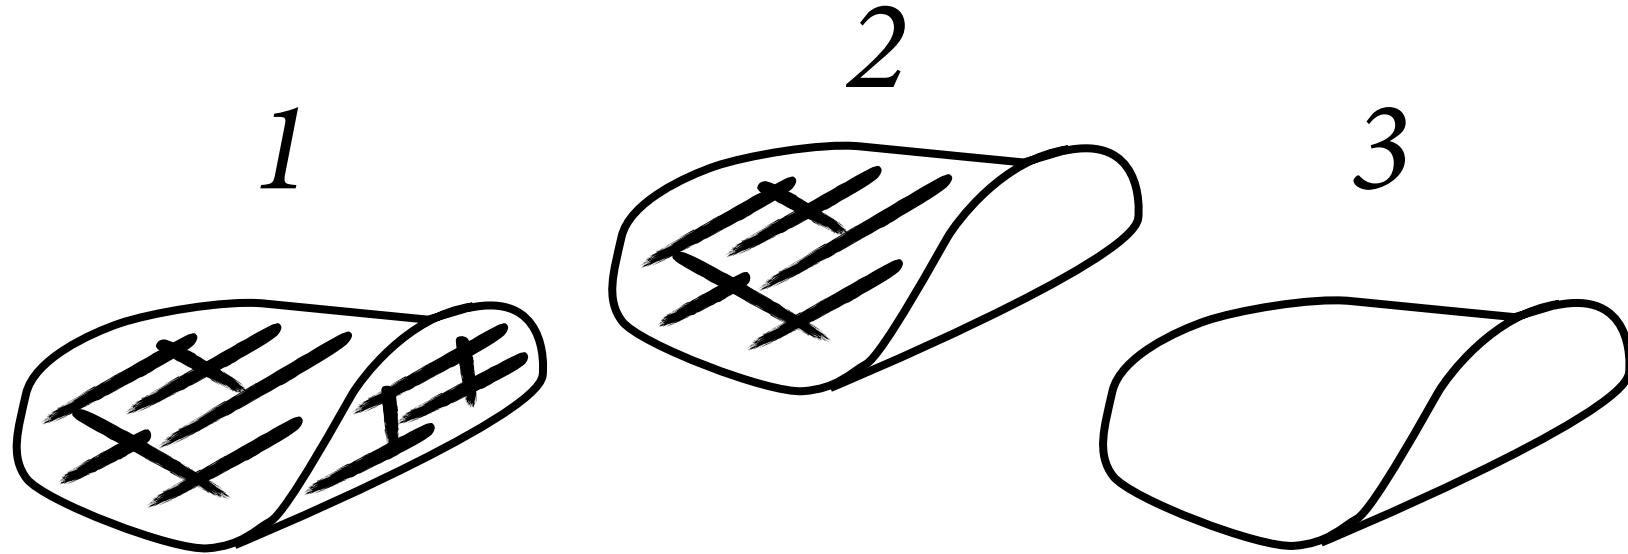
$$\begin{aligned}\Pr(\text{burnt up}) &= \Pr(\text{BB})(1) + \Pr(\text{BU})(0.5) + \Pr(\text{UU})(0) \\ &= (1/3) + (1/3)(1/2) = 0.5\end{aligned}$$



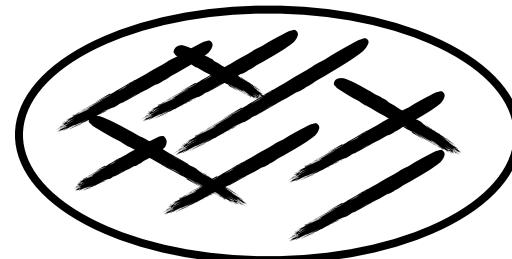
$$\Pr(\text{burnt down} | \text{burnt up}) = \frac{\Pr(\text{burnt up, burnt down})}{\Pr(\text{burnt up})}$$

$$\begin{aligned}\Pr(\text{burnt up}) &= \Pr(\text{BB})(1) + \Pr(\text{BU})(0.5) + \Pr(\text{UU})(0) \\ &= (1/3) + (1/3)(1/2) = 0.5\end{aligned}$$

$$\Pr(\text{burnt down} | \text{burnt up}) = \frac{1/3}{1/2} = \frac{2}{3}$$



You are served:



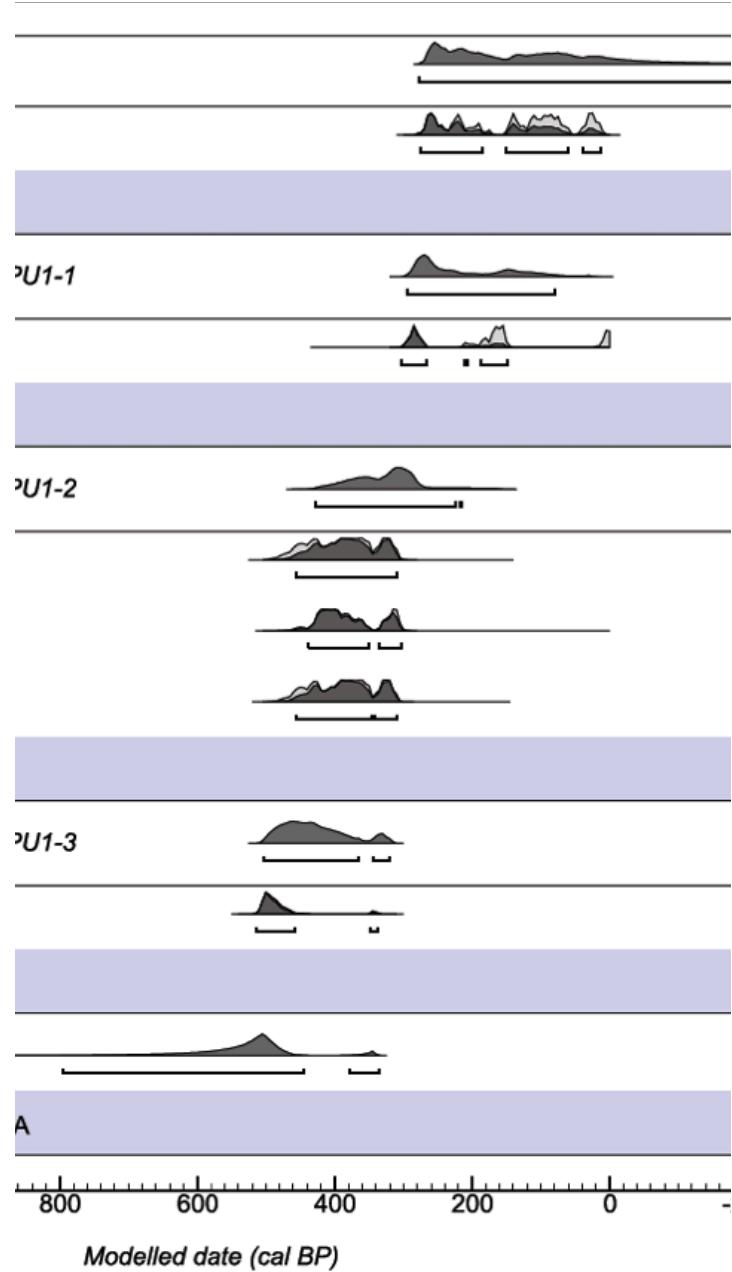
Probability other side is burnt?

Getting Ruthless

- Express information as constraints and distributions => let logic discover implications
- No need to be clever
- Examples:
 - Measurement error
 - Missing data

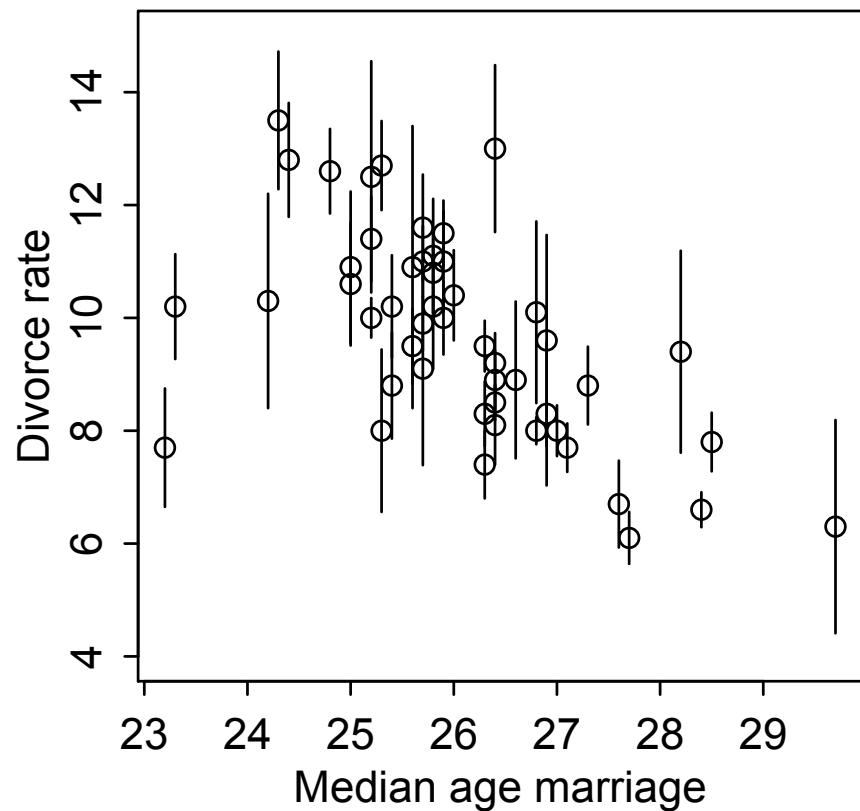
Measurement error

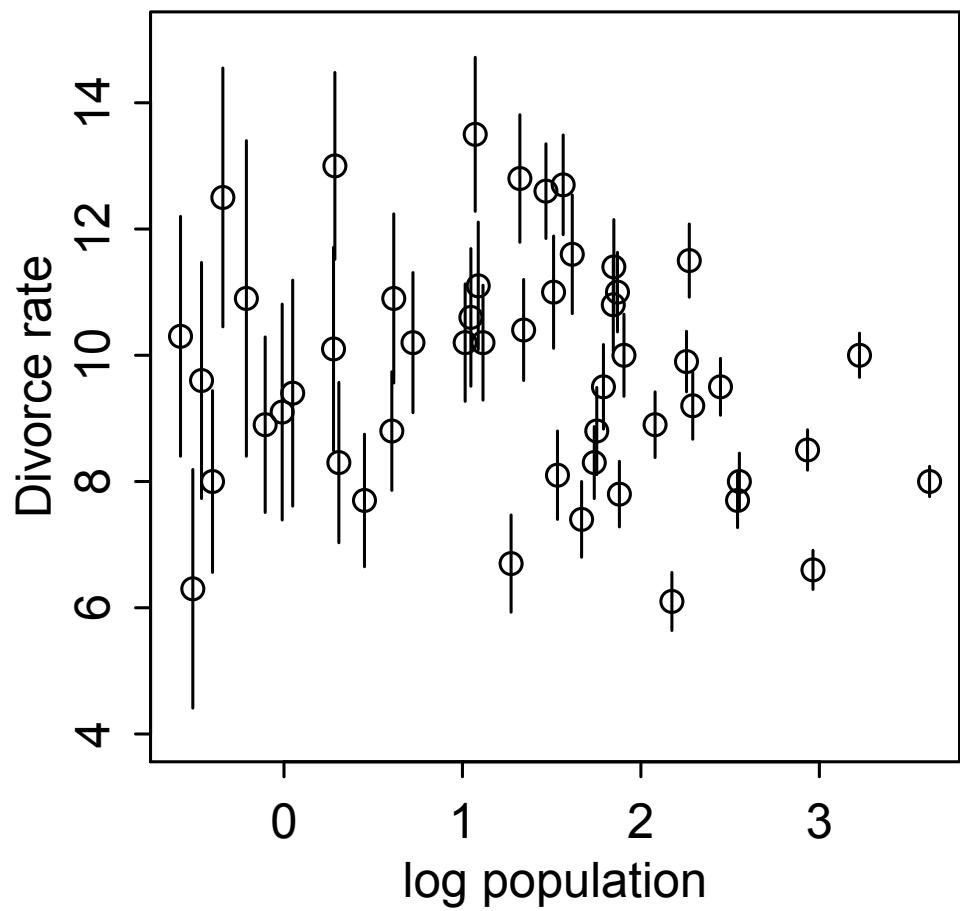
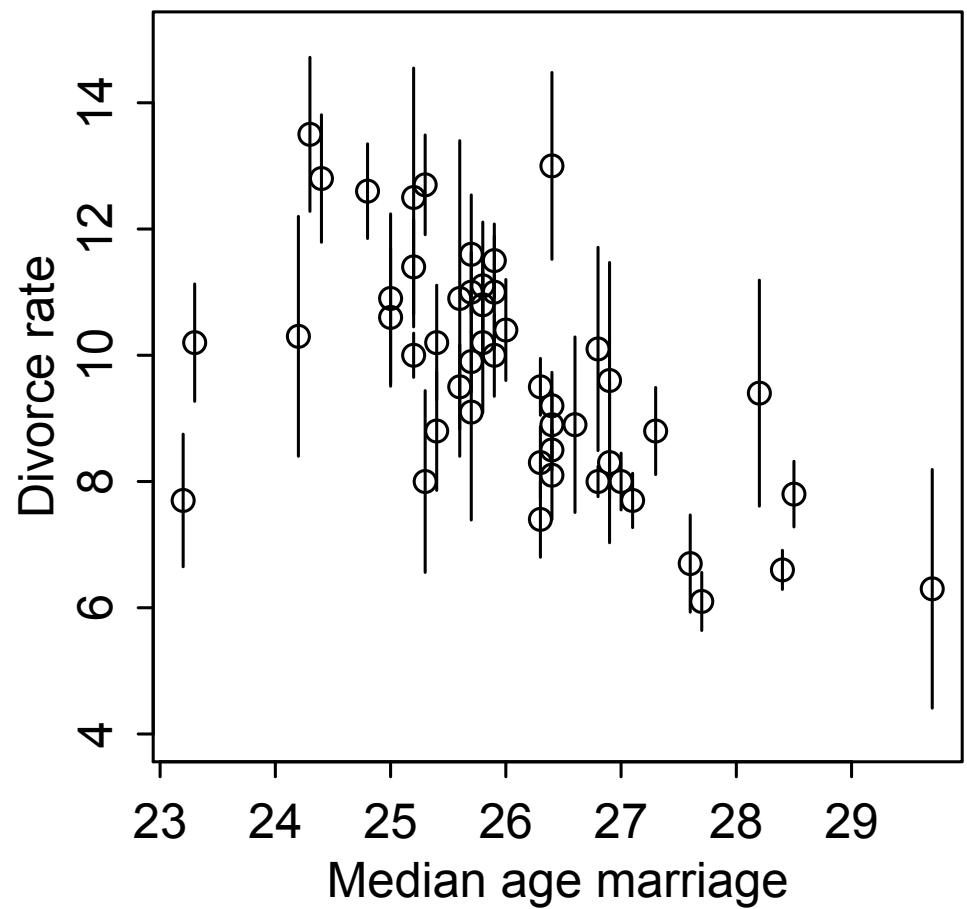
- Measurement always entails error
- Typical linear regression: interpret sigma as “error” on outcome
- What if error isn’t constant?
- What if error is on predictors?

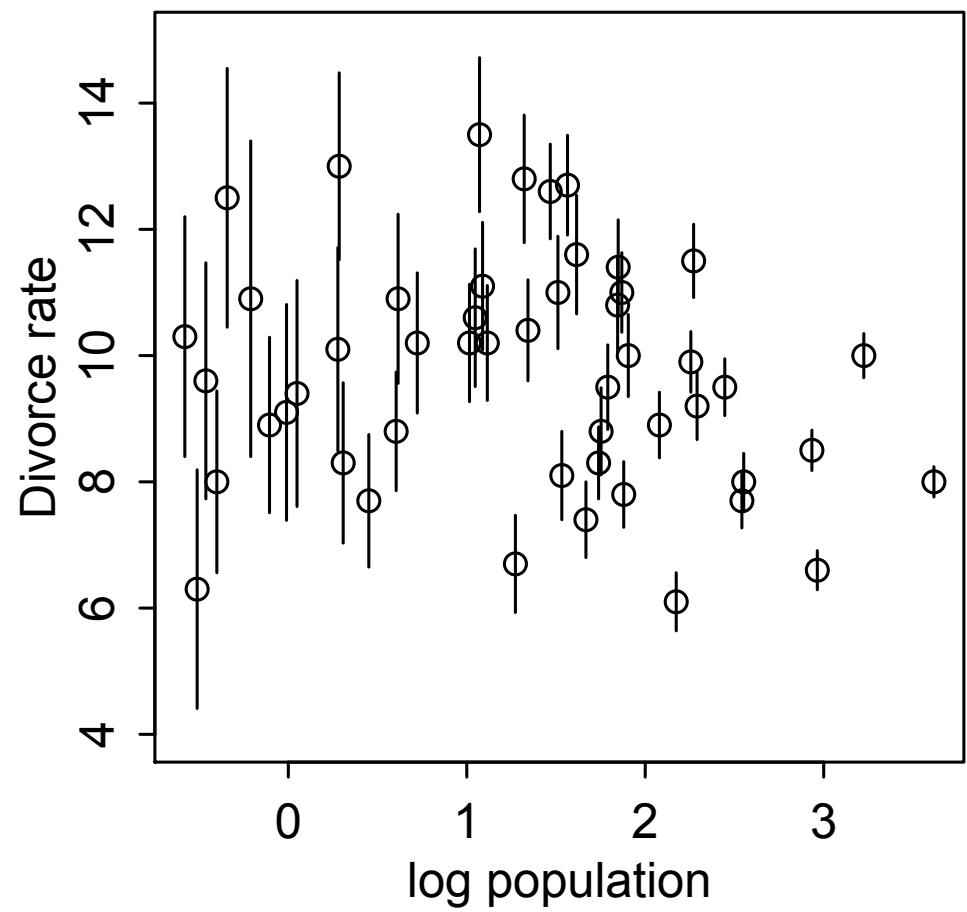
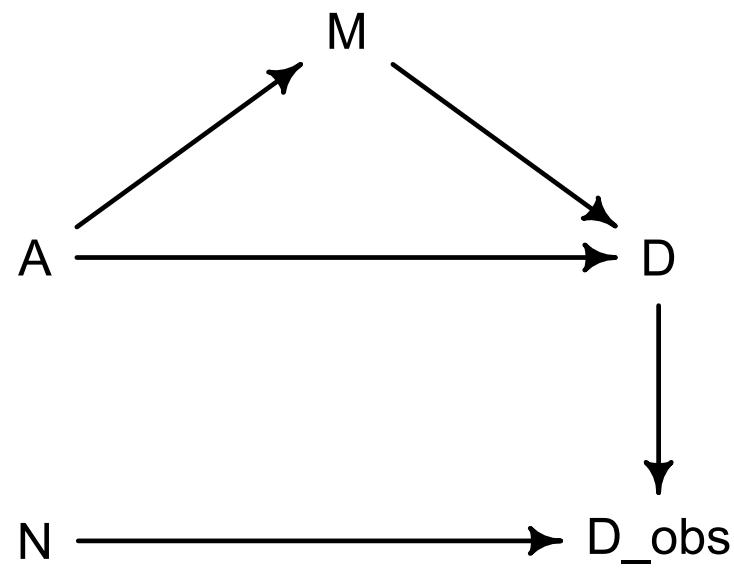


Error on outcome

- `data(WaffleDivorce)`
- Consider error on outcome, divorce rate
- Heterogeneity in error
- Small State => large error







Error on outcome

- Approach:
 - Treat true divorce rate as unknown parameter
 - Observed rate is sample from Gaussian distribution:

$$D_{\text{OBS},i} \sim \text{Normal}(D_{\text{TRUE},i}, D_{\text{SE},i})$$

observed *true* *std error*
(data) *(parameter)* *(data)*

Error on outcome: model

$$D_{\text{TRUE},i} \sim \text{Normal}(\mu_i, \sigma)$$

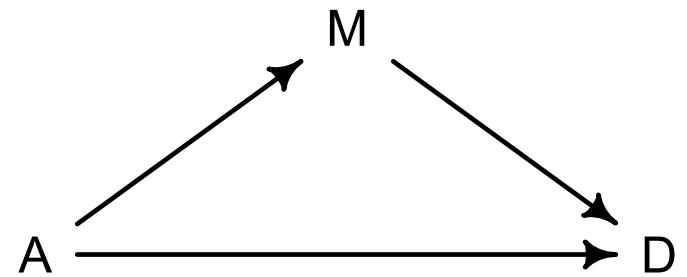
$$\mu_i = \alpha + \beta_A A_i + \beta_M M_i$$

$$\alpha \sim \text{Normal}(0, 0.2)$$

$$\beta_A \sim \text{Normal}(0, 0.5)$$

$$\beta_M \sim \text{Normal}(0, 0.5)$$

$$\sigma \sim \text{Exponential}(1)$$



$$D_{\text{OBS},i} \sim \text{Normal}(D_{\text{TRUE},i}, D_{\text{SE},i})$$

estimate *standard error
of observation*

$$D_{\text{TRUE},i} \sim \text{Normal}(\mu_i, \sigma)$$

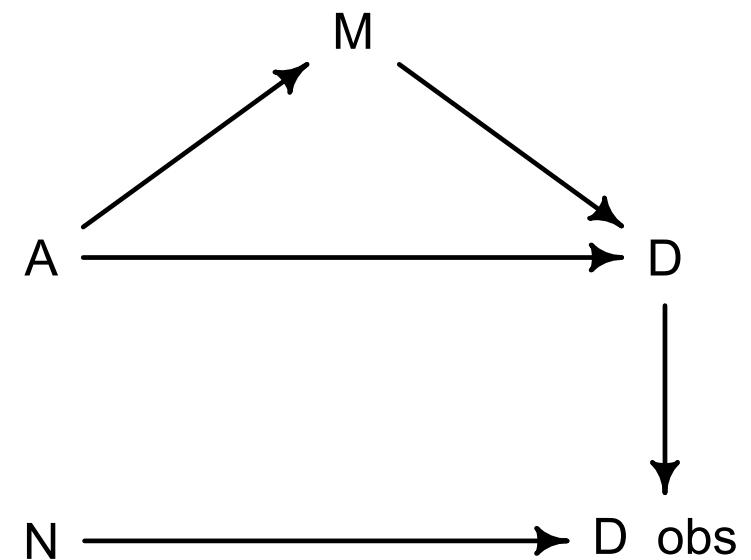
$$\mu_i = \alpha + \beta_A A_i + \beta_M M_i$$

$$\alpha \sim \text{Normal}(0, 0.2)$$

$$\beta_A \sim \text{Normal}(0, 0.5)$$

$$\beta_M \sim \text{Normal}(0, 0.5)$$

$$\sigma \sim \text{Exponential}(1)$$

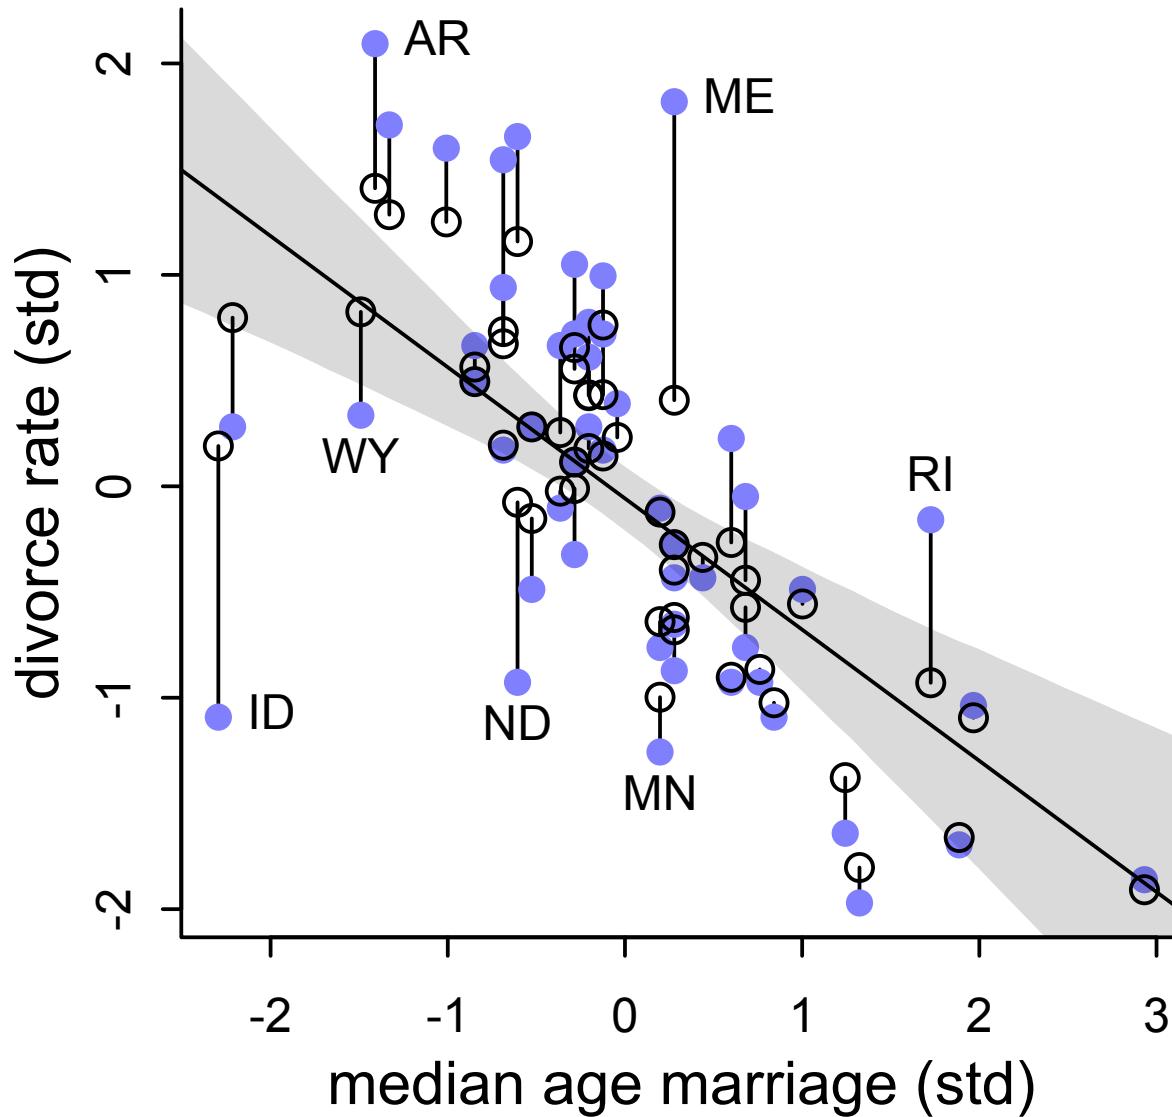


Error on outcome: fitting

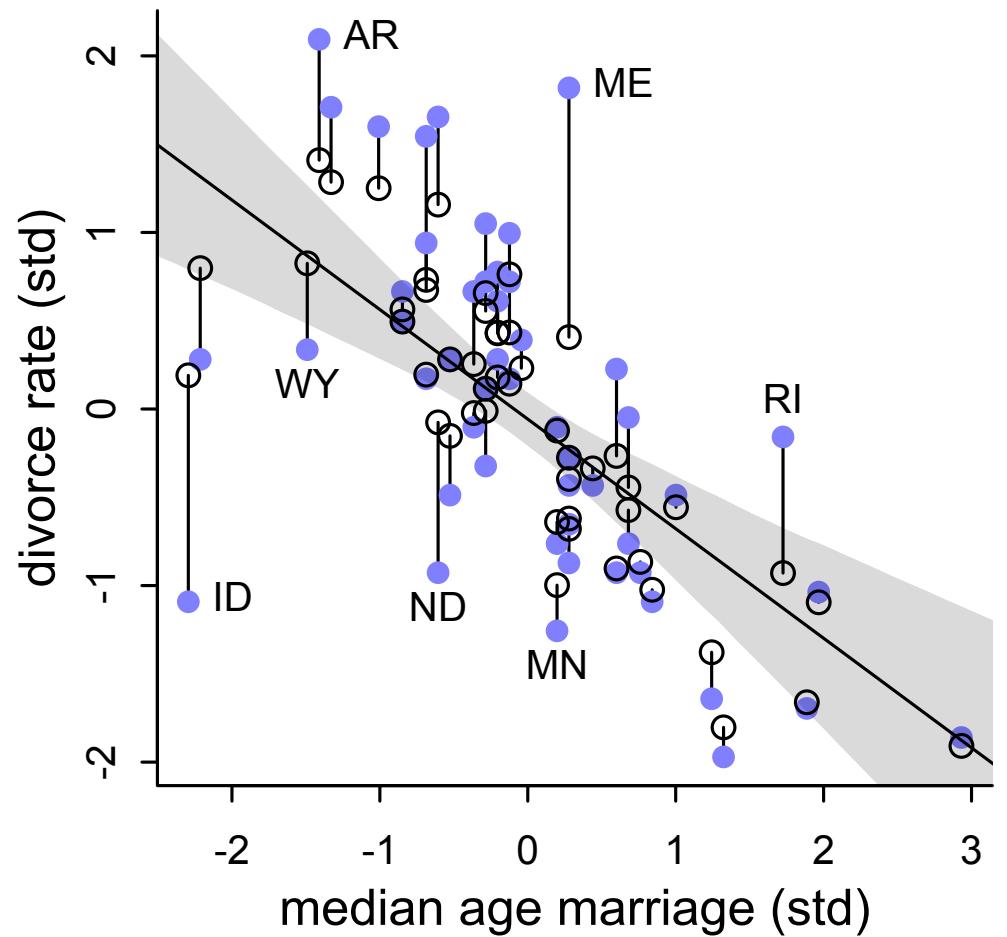
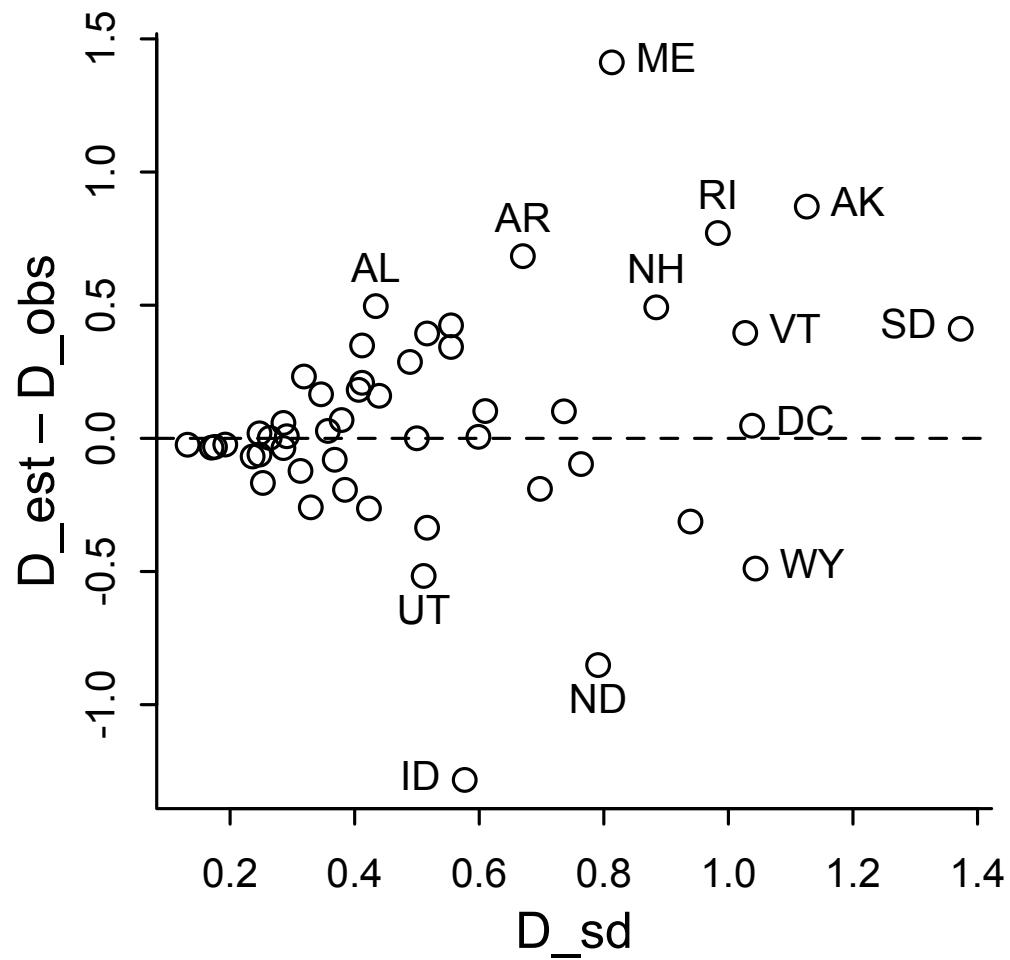
```
m15.1 <- ulam(  
  alist(  
    D_obs ~ dnorm( D_true , D_sd ),  
    vector[N]:D_true ~ dnorm( mu , sigma ),  
    mu <- a + bA*A + bM*M,  
    a ~ dnorm(0,0.2),  
    bA ~ dnorm(0,0.5),  
    bM ~ dnorm(0,0.5),  
    sigma ~ dexp(1)  
  ) , data=dlist , chains=4 , cores=4 )
```

Error on outcome: fitting

```
m15.1 <- ulam(  
  alist(  
    D_obs ~ dnorm( D_true , D_sd ) ,  
    vector[N]:D_true ~ dnorm( mu , sigma ) ,  
    mu <- a + bA*A + bM*M ,  
    a ~ dnorm(0,0.2) ,  
    bA ~ dnorm(0,0.5) ,  
    bM ~ dnorm(0,0.5) ,  
    sigma ~ dexp(1)  
  ) , data=dlist , chains=4 , cores=4 )
```



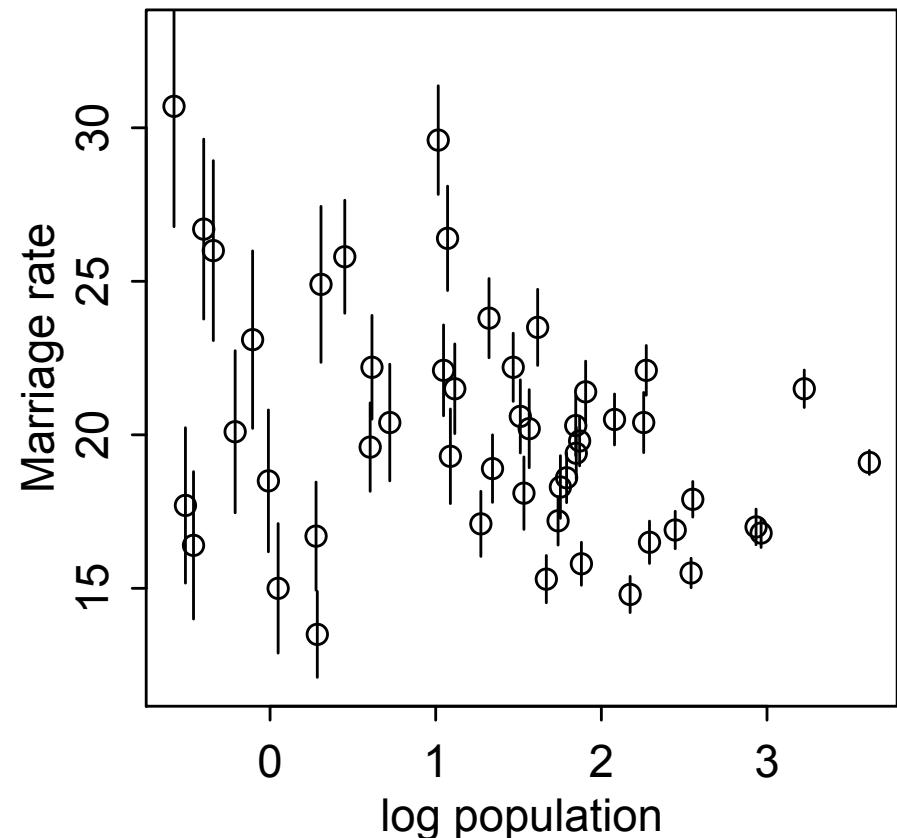
- Divorce rate estimates move from observed values. Why?



- Shrinkage! Uncertain or extreme states shrink to regression line.

Error on predictor

- What about error on predictor?
- Many procedures invented
 - errors-in-variables
 - reduced major axis
 - total least squares
- Our approach will be logical
 - State information
 - Deduce implications
 - Garbage in? You know what comes out.



Error on predictor: model

$$D_{\text{OBS},i} \sim \text{Normal}(D_{\text{TRUE},i}, D_{\text{SE},i})$$

$$D_{\text{TRUE},i} \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta_A A_i + \beta_M M_{\text{TRUE},i}$$

$$M_{\text{OBS},i} \sim \text{Normal}(M_{\text{TRUE},i}, M_{\text{SE},i})$$

$$M_{\text{TRUE},i} \sim \text{Normal}(0, 1)$$

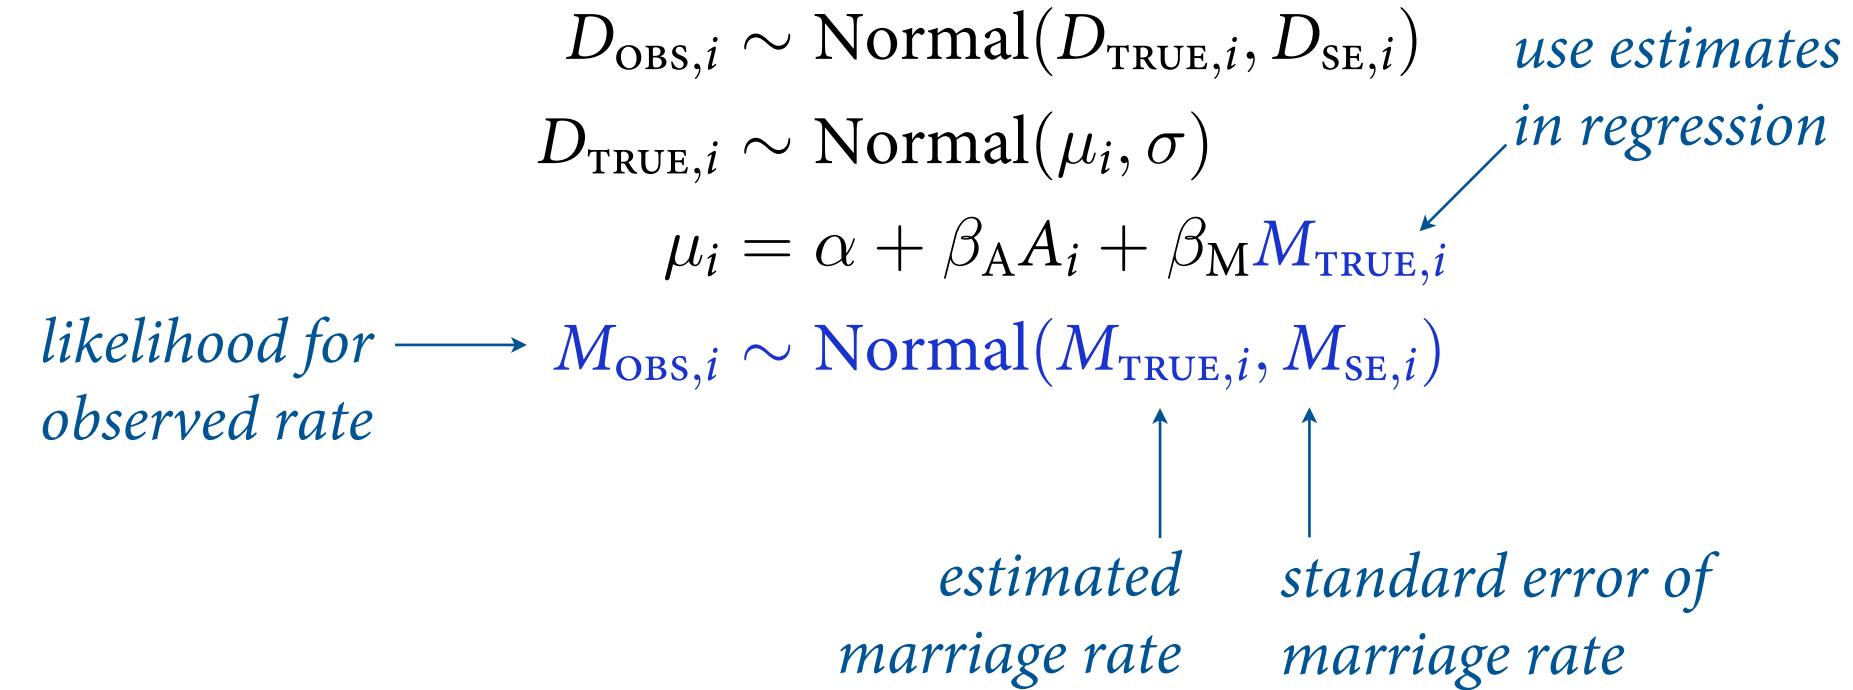
$$\alpha \sim \text{Normal}(0, 0.2)$$

$$\beta_A \sim \text{Normal}(0, 0.5)$$

$$\beta_M \sim \text{Normal}(0, 0.5)$$

$$\sigma \sim \text{Exponential}(1)$$

Error on predictor: model



Error on predictor: model

$$D_{\text{OBS},i} \sim \text{Normal}(D_{\text{TRUE},i}, D_{\text{SE},i})$$

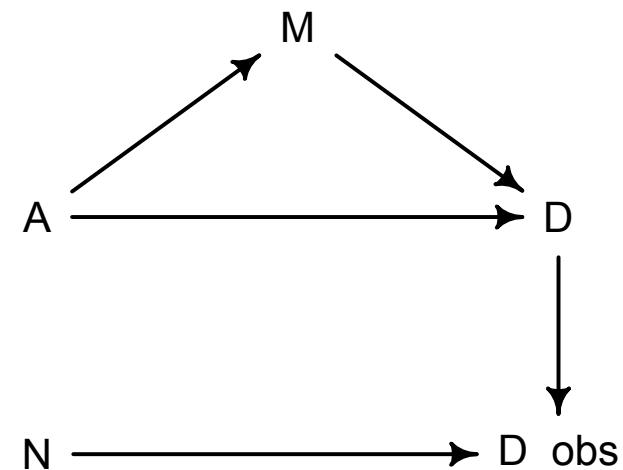
$$D_{\text{TRUE},i} \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta_A A_i + \beta_M M_{\text{TRUE},i}$$

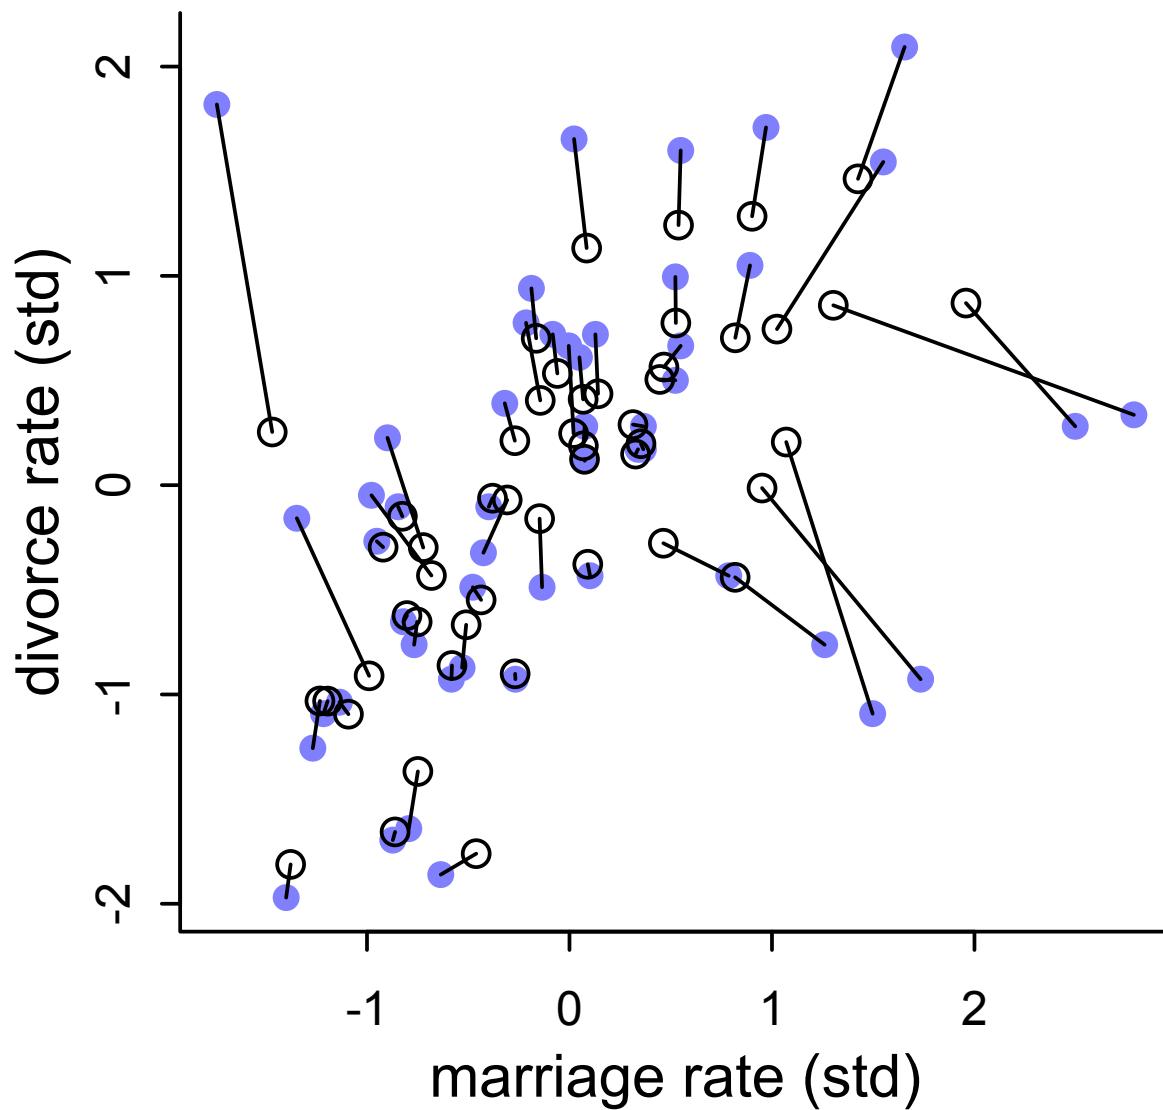
$$M_{\text{OBS},i} \sim \text{Normal}(M_{\text{TRUE},i}, M_{\text{SE},i})$$

prior rates —— $M_{\text{TRUE},i} \sim \text{Normal}(0, 1)$

Not the best approach:
M and A are associated!
Will do better later on.

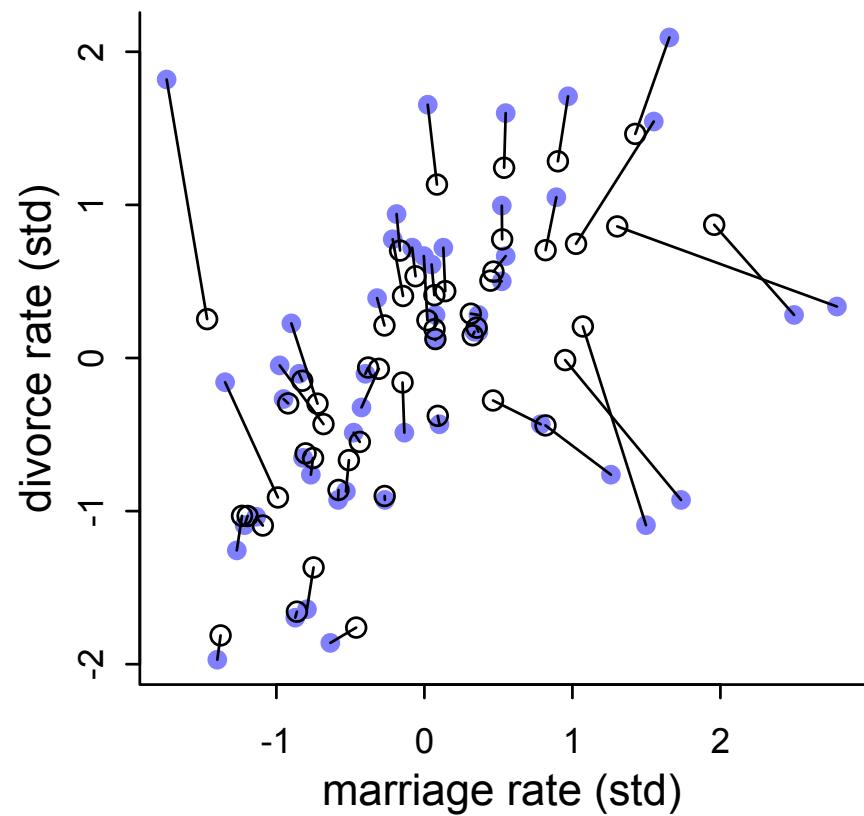


filled circles: observed
open circles: estimated
lines connect points for same State



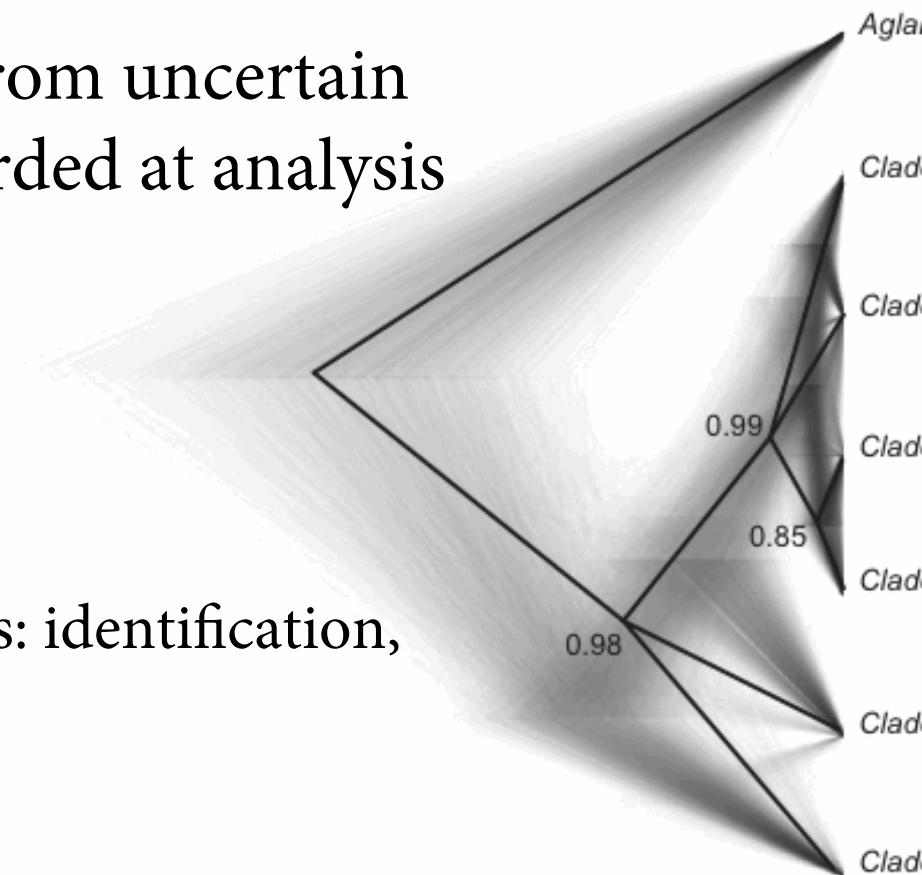
Error on predictor

- Both divorce rate and marriage rate shrink
- Divorce shrinks more. Why?
- Marriage rate not strongly associated with outcome => not much pooling through regression => not much shrinkage



Measurement error

- Common malady: “data” come from uncertain procedure, but uncertainty discarded at analysis
- Examples:
 - Predicting with averages
 - Parentage analysis
 - Phylogenetics: distribution of trees
 - Archaeology/paleontology/forensics: identification, sexing, aging, dating
- Propagate uncertainty

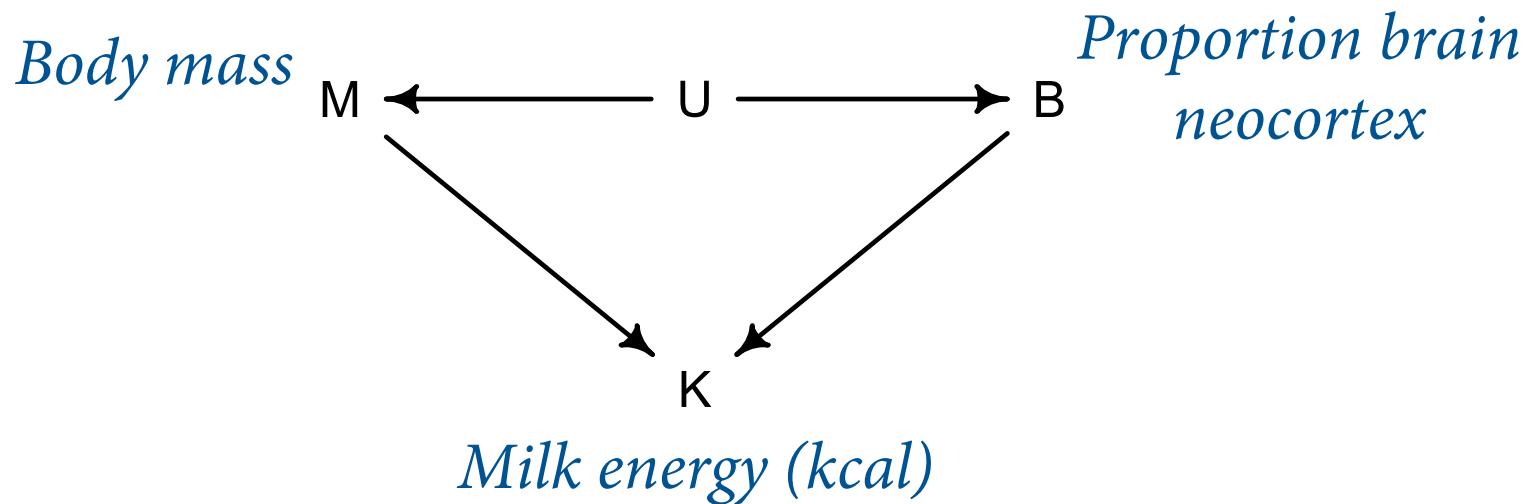


Missing data

- Missing values commonplace
 - Usual approach: **complete-case** analysis
 - drop all cases with any missing values
 - Discards a lot of information
 - Alternatives
 - replace missing with mean of column: NEVER DO THIS
 - Multiple imputation
 - Bayesian imputation
 - others
- im•pute** | im'pyoot |
verb [with obj.]
represent (something, esp. something undesirable) as being done,
caused, or possessed by someone; attribute: *the crimes imputed
to Richard*.
- Finance assign (a value) to something by inference from the value
of the products or processes to which it contributes: (as adj.
imputed) : *recovering the initial outlay plus imputed interest.*
 - Theology ascribe (righteousness, guilt, etc.) to someone by virtue of
a similar quality in another: *Christ's righteousness has been imputed
to us.*

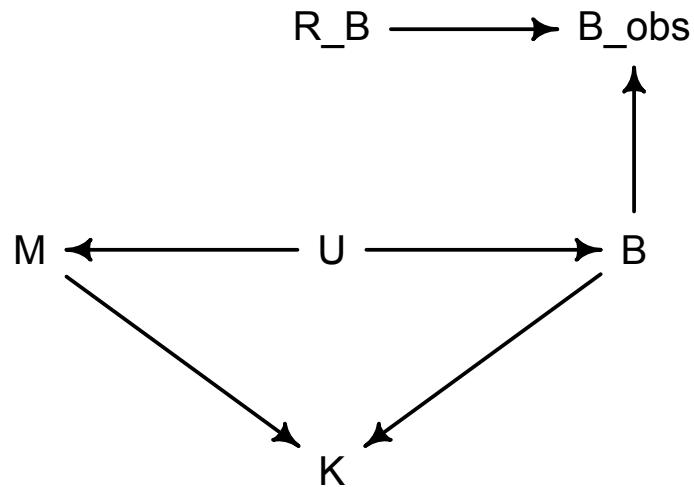
Why impute?

- Missingness can be a confound

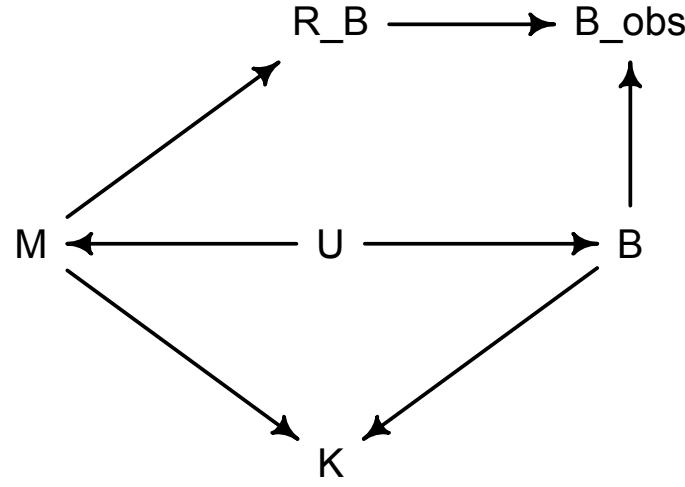


Three Types of Missingness

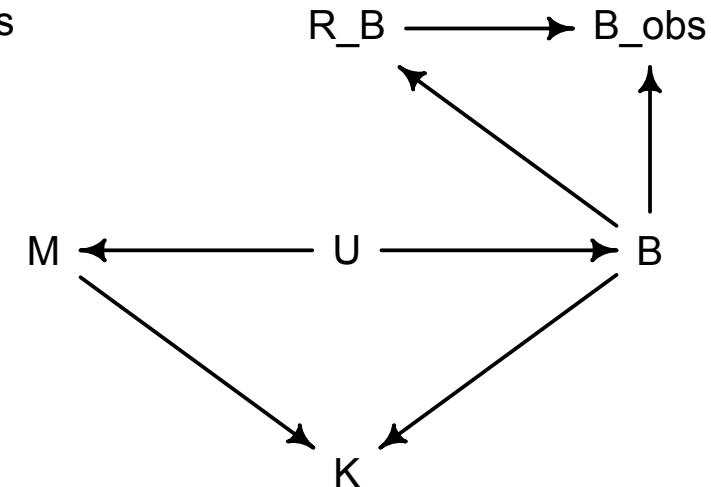
MCAR



MAR



MNAR



MISSING COMPLETELY
AT RANDOM

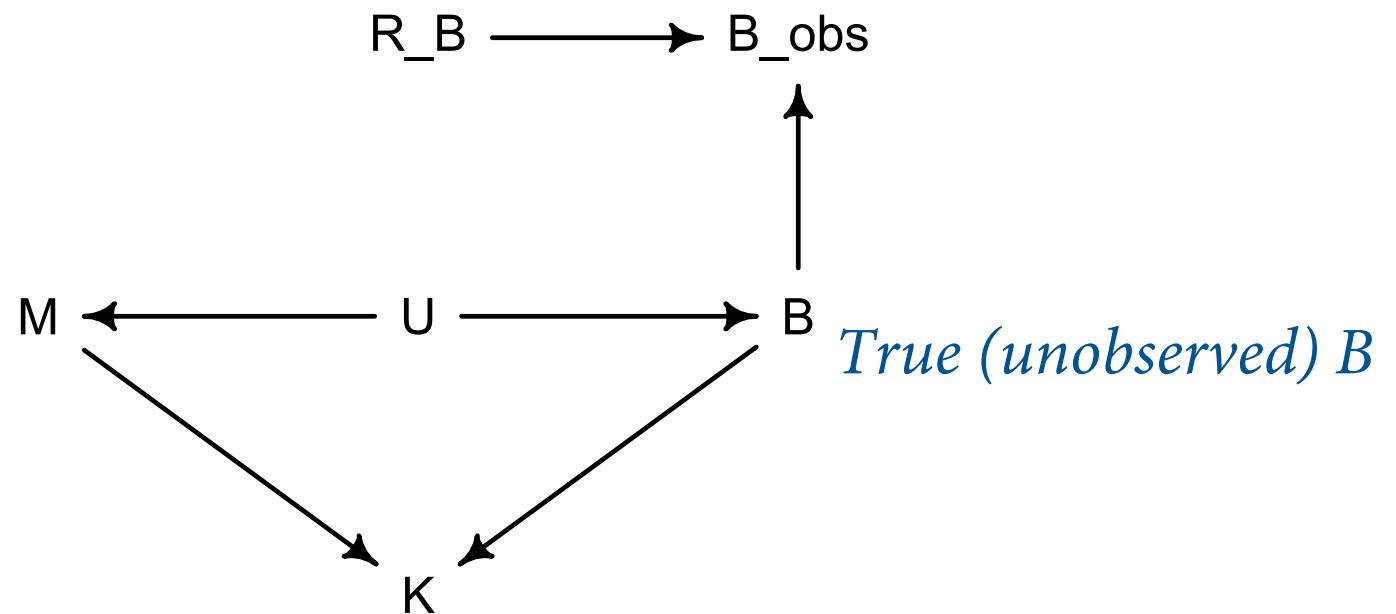
MISSING
AT RANDOM

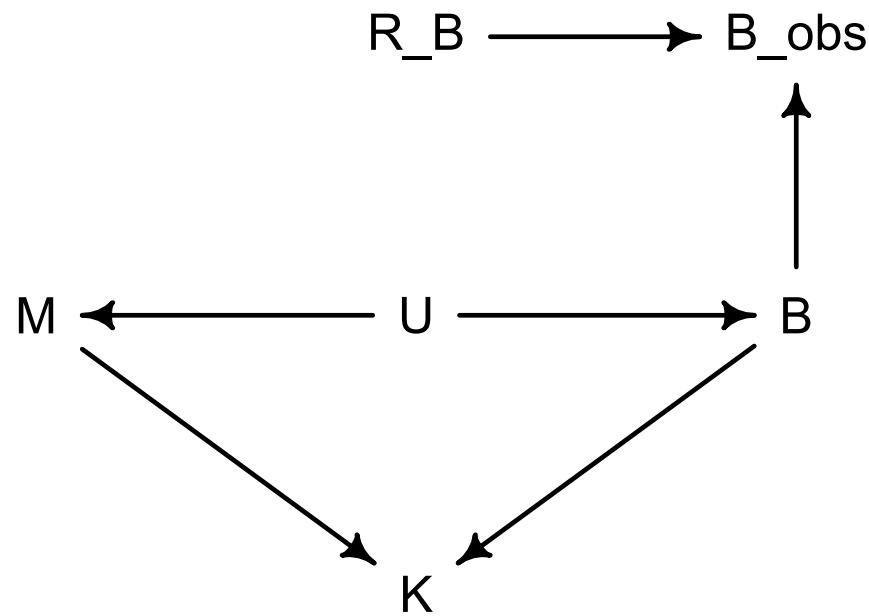
MISSING NOT
AT RANDOM

Possibly most confusing statistical terms ever invented.

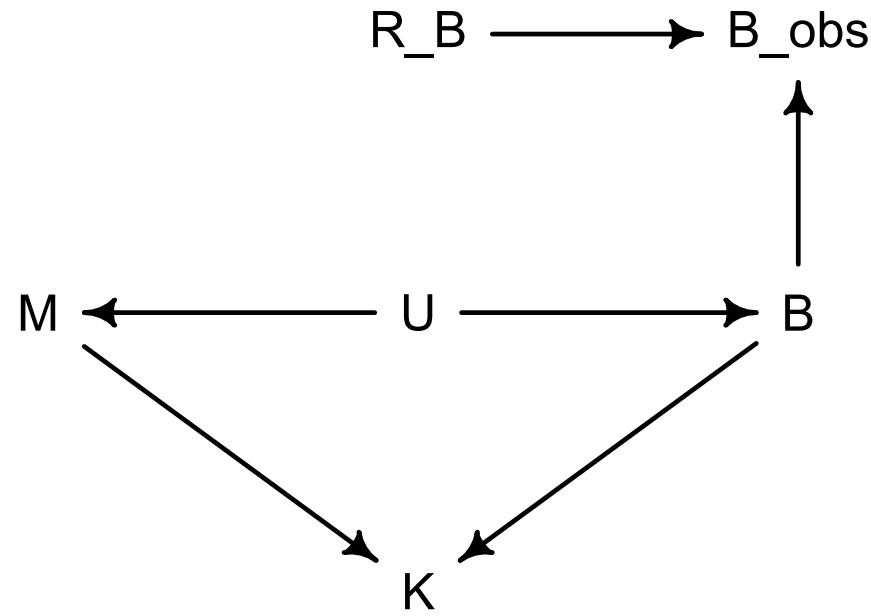
*Missingness
mechanism*

Observed B





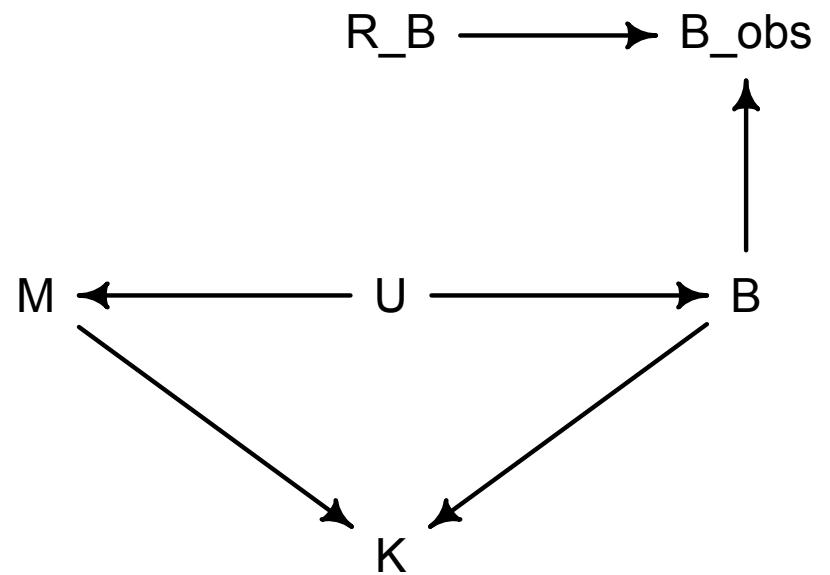
Are there any backdoors from B_{obs} to K ?

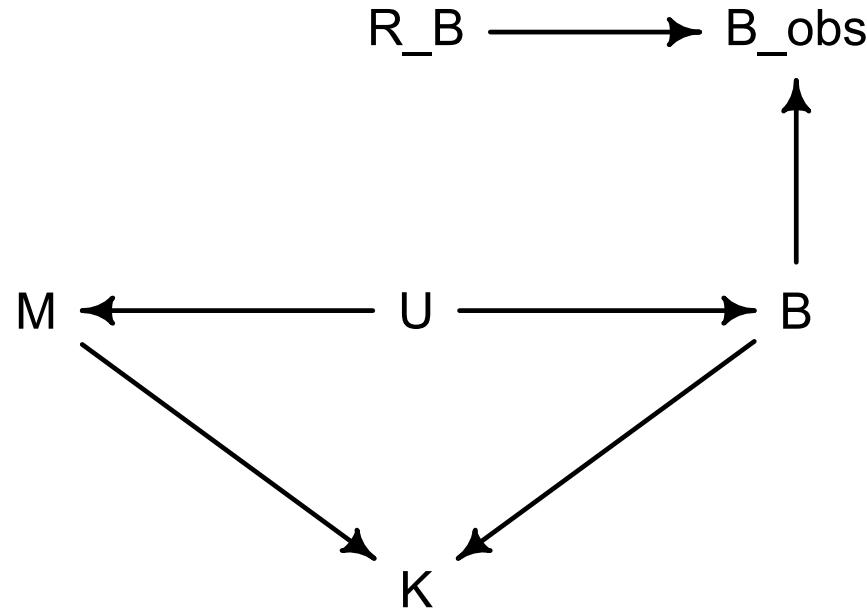


Can condition on M for direct effect.
Either way, R_B is ignorable.

Missing Completely At Random

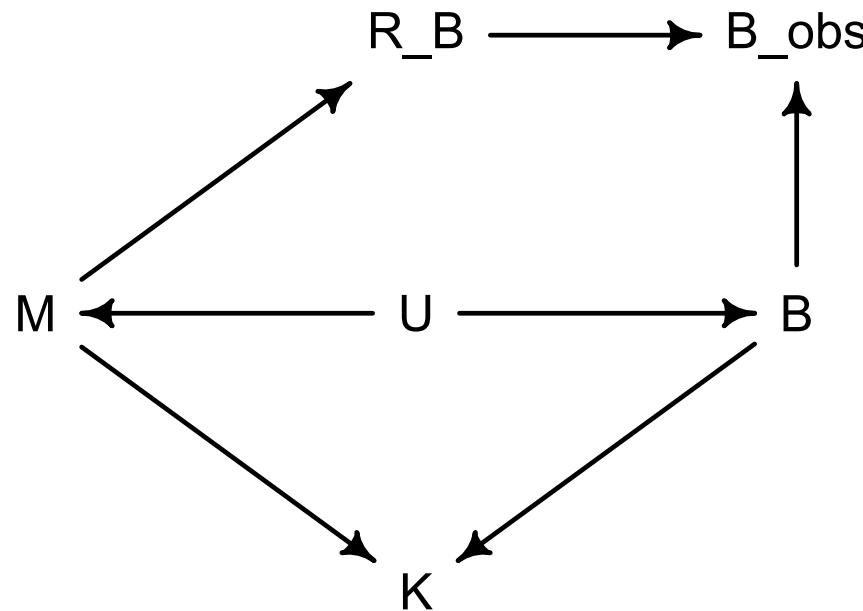
- MCAR: K is unconditionally independent of R_B
- Do not need to condition on anything for R_B not to be a confound
- On right, no path through R_B, conditioning on B_obs
- Do not NEED to impute
- But imputation adds precision





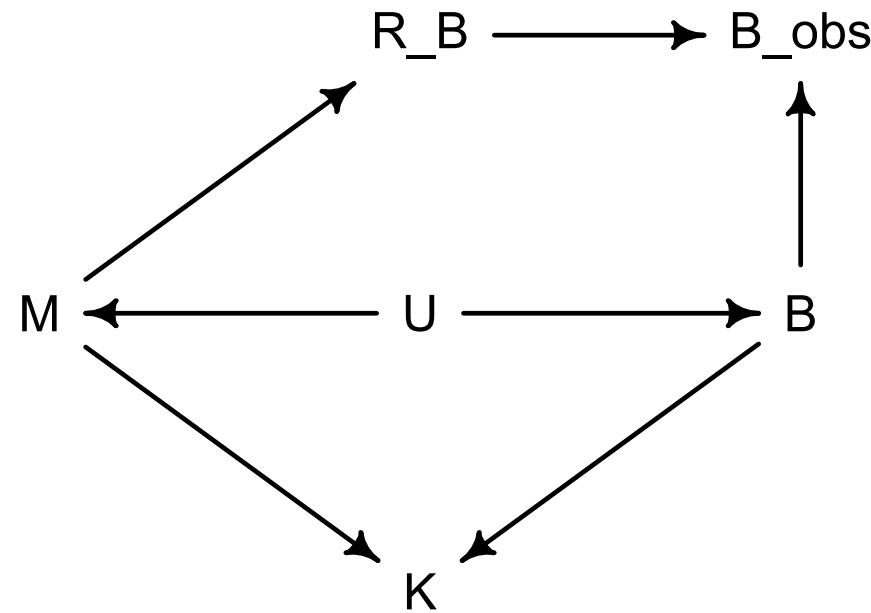
Does MCAR ever happen in real data?
Research assistant randomly deletes values?

Missing At Random



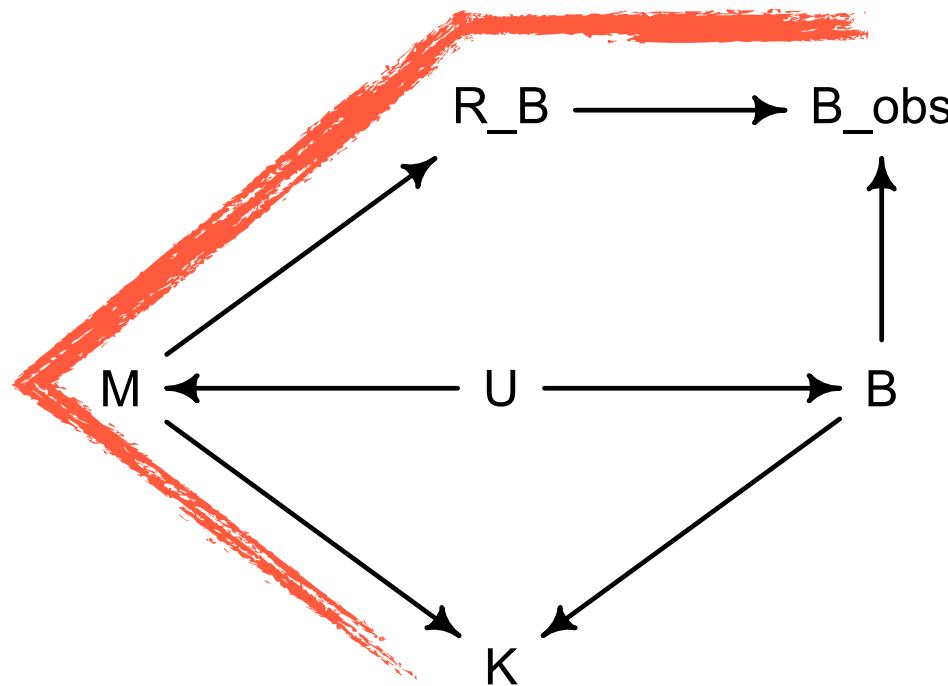
Missingness more likely for specific values of M .
How can this happen?

Missing At Random



Backdoor path from B_{obs} to K ?

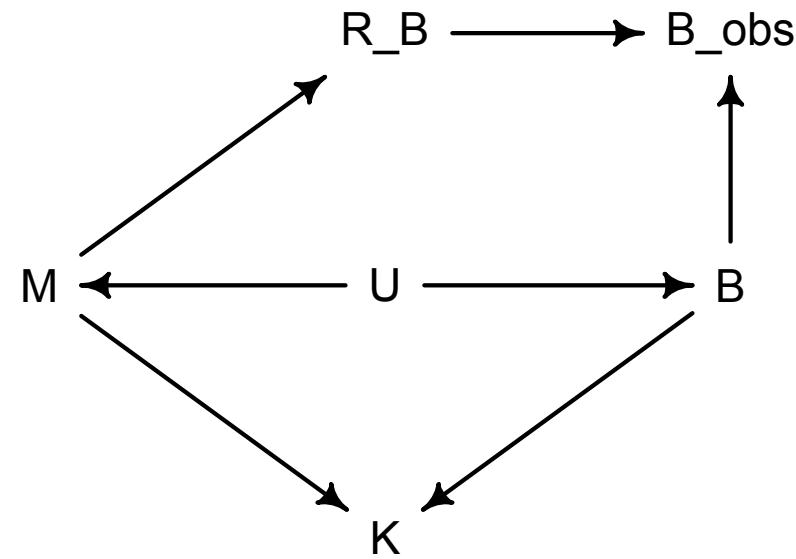
Missing At Random



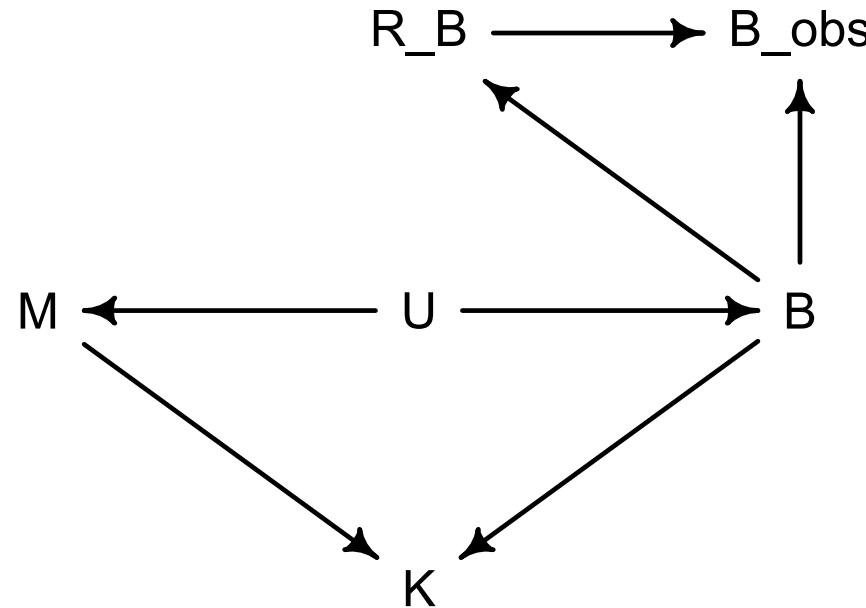
Backdoor path from B_{obs} to K ?
Can condition on M to close.

Missing (Simply) At Random

- MAR: K is conditionally independent of R_B
- Must to condition on M for R_B not to be a confound
- Still must impute to de-bias estimates
- Why? If you delete cases of M/K where B is missing, missingness obscures causation.

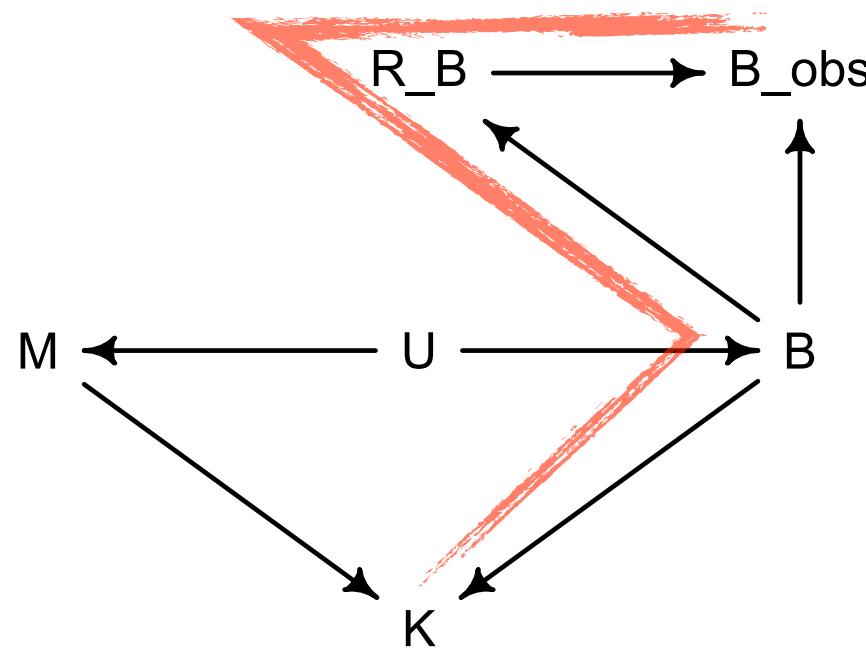


Missing Not At Random



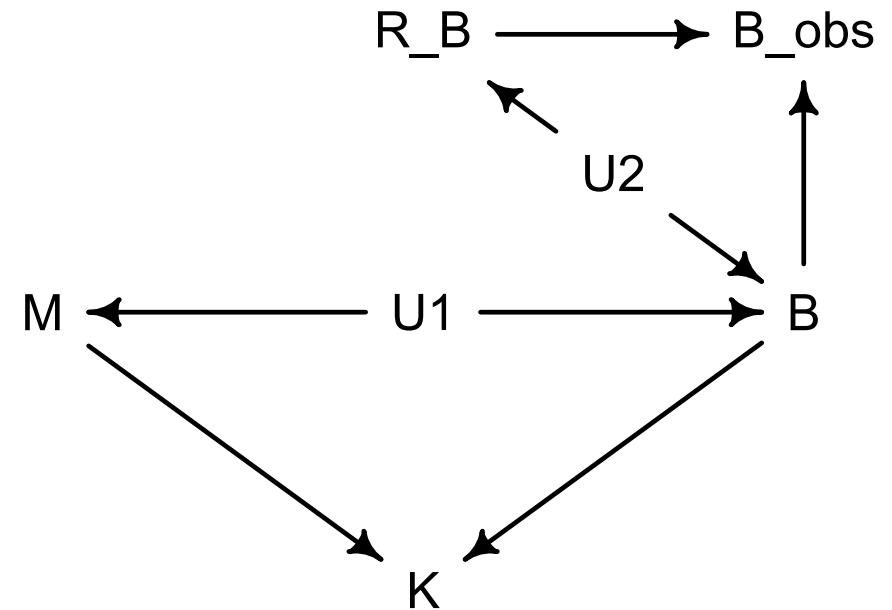
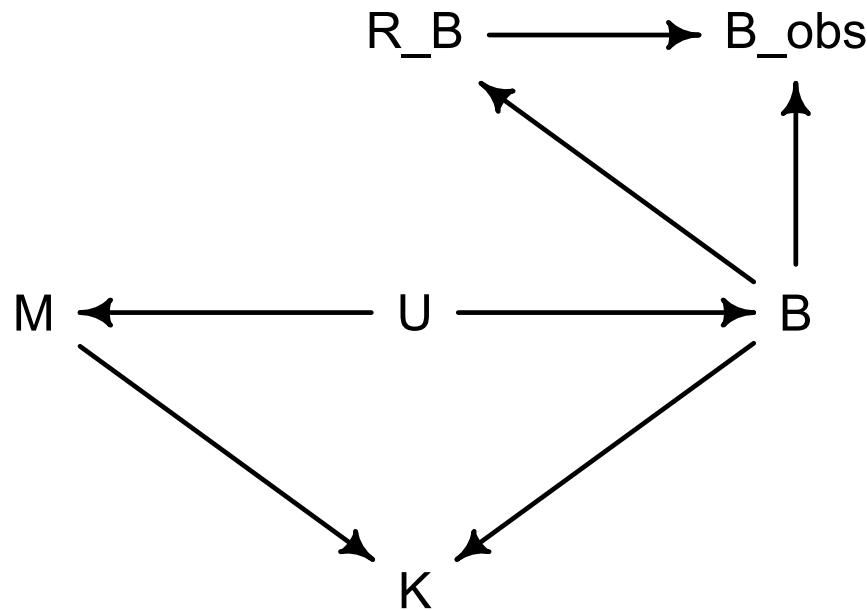
Missingness more likely for specific values of B.
How can this happen?

Missing Not At Random



No way to shut the backdoor!

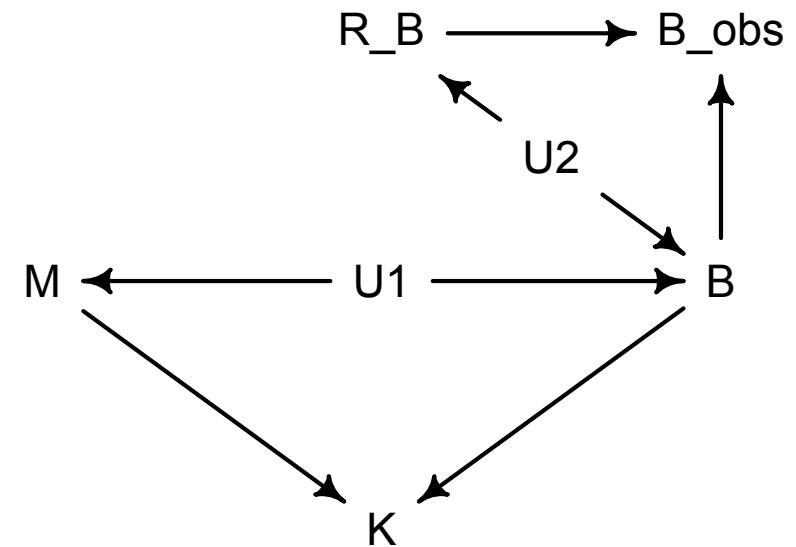
Missing Not At Random



Can also arise through unobserved variables (right).

Missing Not At Random

- MNAR: K is unconditionally **dependent** on R_B
- If you can model R_B , might be okay
- No guarantees



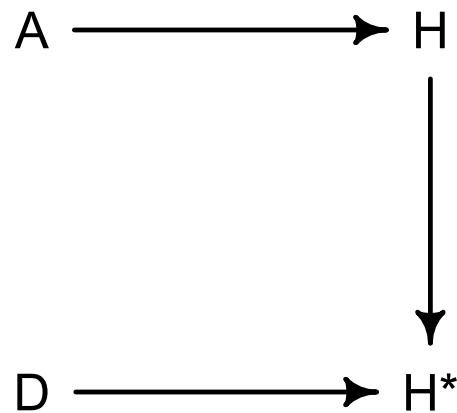
H: Homework

H*: Homework with missing values

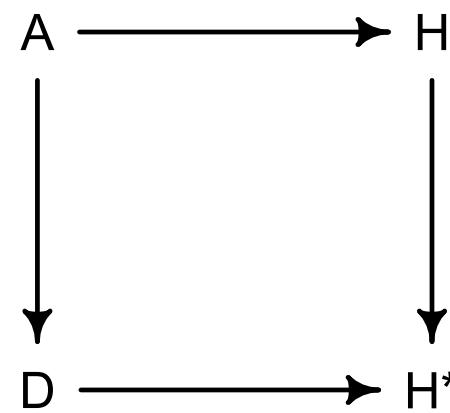
A: Attribute of student

D: Dog (missingness mechanism)

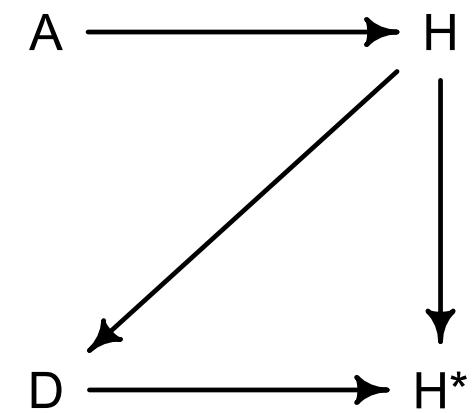
DOG EATS
ANY
HOMEWORK



DOG EATS
STUDENTS'
HOMEWORK



DOG EATS
BAD
HOMEWORK



MISSING COMPLETELY
AT RANDOM

MISSING
AT RANDOM

MISSING NOT
AT RANDOM

Milk imputation

- `data(milk)`
- 12 missing values for neocortex
- Suppose values are *Missing At Random (MAR)*
 - Distribution of observed values provides information
 - Can use to impute missing values
 - Same procedure for MCAR

	kcal.per.g	mass	neocortex.perc
1	0.49	1.95	55.16
2	0.51	2.09	NA
3	0.46	2.51	NA
4	0.48	1.62	NA
5	0.60	2.19	NA
6	0.47	5.25	64.54
7	0.56	5.37	64.54
8	0.89	2.51	67.64
9	0.91	0.71	NA
10	0.92	0.68	68.85
11	0.80	0.12	58.85
12	0.46	0.47	61.69
13	0.71	0.32	60.32
14	0.71	0.60	NA
15	0.73	3.47	NA
16	0.68	1.55	69.97
17	0.72	7.08	NA
18	0.97	3.24	70.41
19	0.79	7.94	NA
20	0.84	12.30	73.40
21	0.48	7.59	NA
22	0.62	5.37	67.53
23	0.51	10.72	NA
24	0.54	35.48	71.26
25	0.49	79.43	72.60
26	0.53	97.72	NA
27	0.48	40.74	70.24
28	0.55	33.11	76.30
29	0.71	54.95	75.49

Milk energy MAR

- Consider just neocortex variable:
 - Q: What is your best guess of each missing value?
 - A: Posterior distribution derived from remaining data

	neocortex.perc
1	55.16
2	NA
3	NA
4	NA
5	NA
6	64.54
7	64.54
8	67.64
9	NA
10	68.85
11	58.85
12	61.69
13	60.32
14	NA
15	NA
16	69.97
17	NA
18	70.41
19	NA
20	73.40
21	NA
22	67.53
23	NA
24	71.26
25	72.60
26	NA
27	70.24
28	76.30
29	75.49

Milk energy MCAR

- Place a unique parameter for each missing value
 - NC1 ... NC12
 - These are values to be imputed

	neocortex.perc
1	55.16
2	NC1
3	NC2
4	NC3
5	NC4
6	64.54
7	64.54
8	67.64
9	NC5
10	68.85
11	58.85
12	61.69
13	60.32
14	NC6
15	NC7
16	69.97
17	NC8
18	70.41
19	NC9
20	73.40
21	NC10
22	67.53
23	NC11
24	71.26
25	72.60
26	NC12
27	70.24
28	76.30
29	75.49

Milk energy MAR: model

$$B = [0.55, B_2, B_3, B_4, 0.65, 0.65, \dots, 0.76, 0.75]$$

$$K_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta_B B_i + \beta_M \log M_i$$

$$B_i \sim \text{Normal}(\nu, \sigma_B)$$

$$\alpha \sim \text{Normal}(0, 0.5)$$

$$\beta_B \sim \text{Normal}(0, 0.5)$$

$$\beta_M \sim \text{Normal}(0, 0.5)$$

$$\sigma \sim \text{Exponential}(1)$$

$$\nu \sim \text{Normal}(0.5, 1)$$

$$\sigma_B \sim \text{Exponential}(1)$$

Milk energy MAR: model

$$B = [0.55, B_2, B_3, B_4, 0.65, 0.65, \dots, 0.76, 0.75]$$

*linear model using
mix of observed and imputed values* $\longrightarrow \mu_i = \alpha + \beta_B \boxed{B_i} + \beta_M \log M_i$

$$K_i \sim \text{Normal}(\mu_i, \sigma)$$
$$B_i \sim \text{Normal}(\nu, \sigma_B)$$
$$\alpha \sim \text{Normal}(0, 0.5)$$
$$\beta_B \sim \text{Normal}(0, 0.5)$$
$$\beta_M \sim \text{Normal}(0, 0.5)$$
$$\sigma \sim \text{Exponential}(1)$$
$$\nu \sim \text{Normal}(0.5, 1)$$
$$\sigma_B \sim \text{Exponential}(1)$$

Milk energy MAR: model

$$B = [0.55, B_2, B_3, B_4, 0.65, 0.65, \dots, 0.76, 0.75]$$

$$K_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta_B B_i + \beta_M \log M_i$$

$$B_i \sim \text{Normal}(\nu, \sigma_B)$$

*when obs, likelihood;
when imputed, prior*

*mean neocortex
(to be estimated)*

*std dev of neocortex
(to be estimated)*

Fitting

```
m15.3 <- ulam(  
  alist(  
    K ~ dnorm( mu , sigma ),  
    mu <- a + bB*B + bM*M,  
    B ~ dnorm( nu , sigma_B ),  
    c(a,nu) ~ dnorm( 0 , 0.5 ),  
    c(bB,bM) ~ dnorm( 0, 0.5 ),  
    sigma_B ~ dexp( 1 ),  
    sigma ~ dexp( 1 )  
  ) , data=dat_list , chains=4 , cores=4 )
```

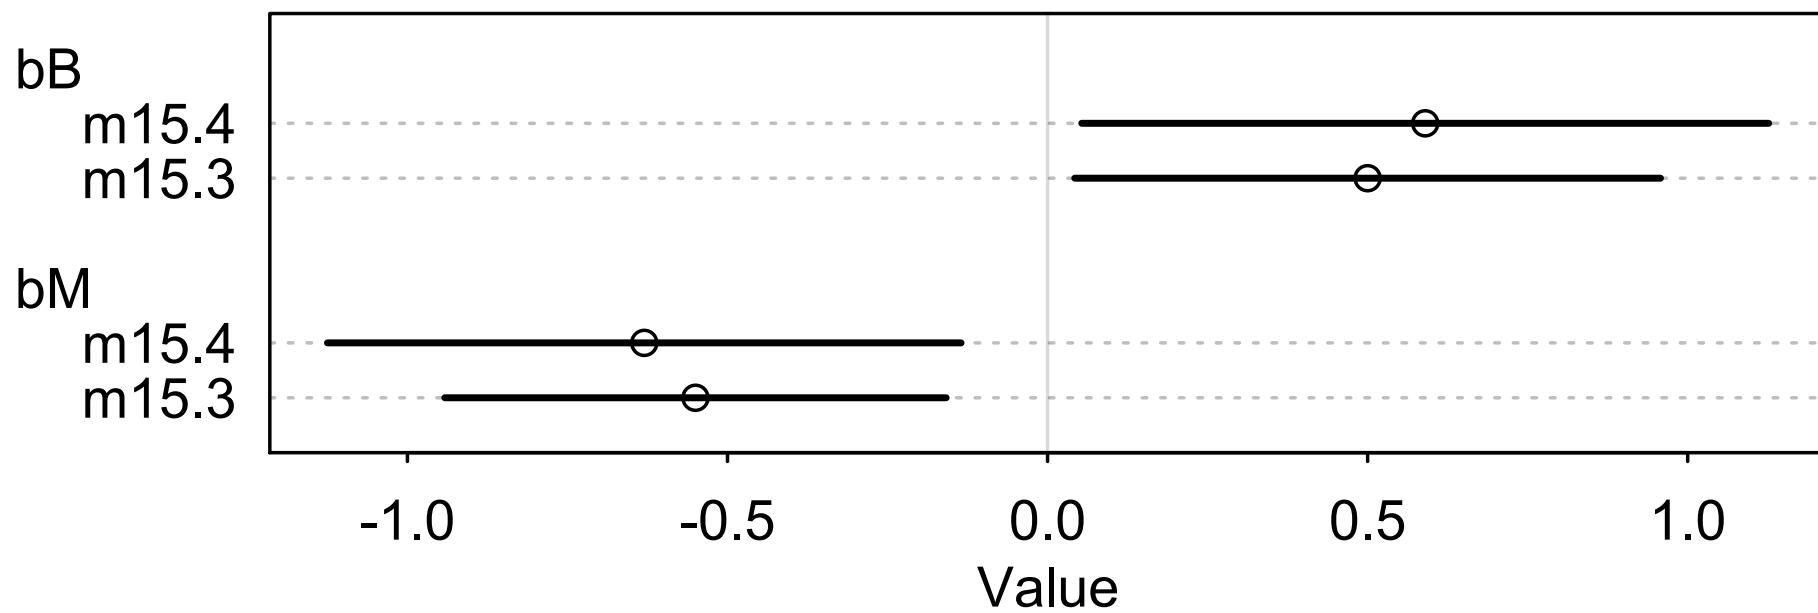
ulam detects NA values and tries to cope.
More explicit example in text.

R code
15.9

```
precis( m15.3 , depth=2 )
```

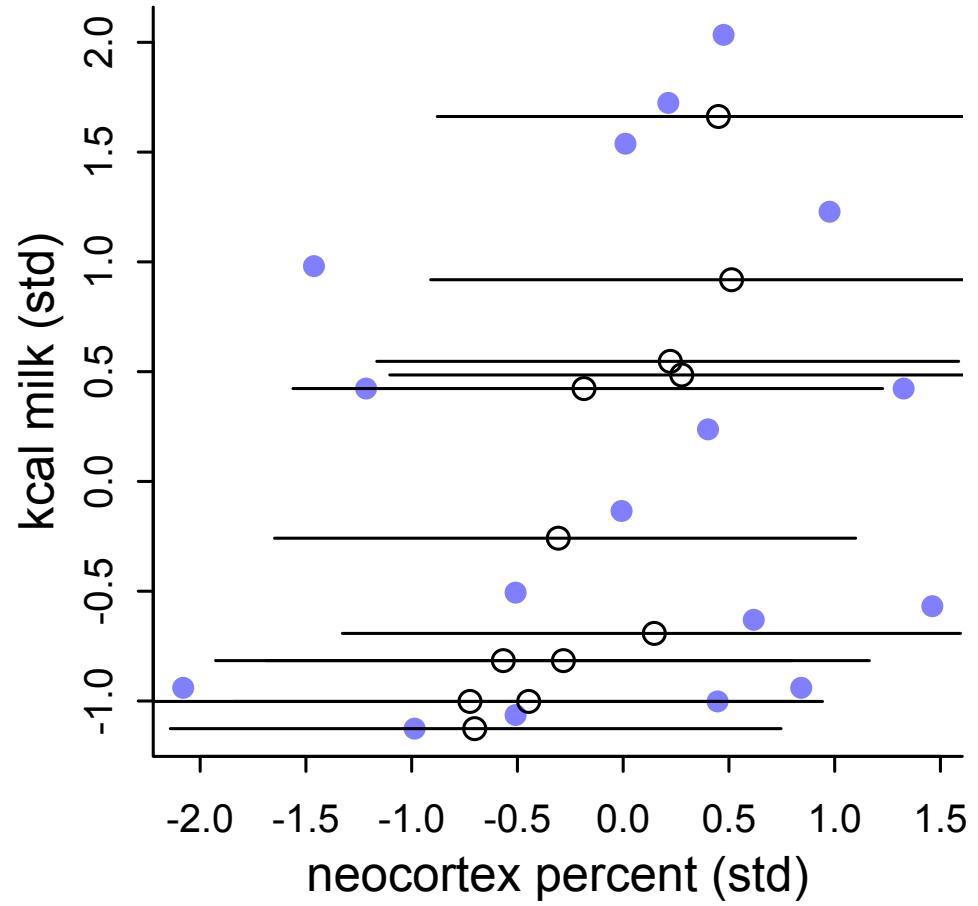
	mean	sd	5.5%	94.5%	n_eff	Rhat
nu	-0.04	0.20	-0.36	0.28	1806	1
a	0.02	0.17	-0.25	0.28	2080	1
bM	-0.55	0.20	-0.87	-0.22	1084	1
bB	0.50	0.23	0.12	0.86	818	1
sigma_B	1.01	0.17	0.77	1.31	1245	1
sigma	0.84	0.14	0.64	1.07	1010	1
B_impute[1]	-0.57	0.87	-1.93	0.80	2408	1
B_impute[2]	-0.70	0.92	-2.14	0.75	1940	1
B_impute[3]	-0.72	0.96	-2.24	0.81	1832	1
B_impute[4]	-0.31	0.87	-1.65	1.10	2345	1
B_impute[5]	0.45	0.88	-0.88	1.83	2416	1
B_impute[6]	-0.19	0.89	-1.56	1.23	2187	1
B_impute[7]	0.22	0.87	-1.16	1.59	2428	1
B_impute[8]	0.28	0.85	-1.10	1.62	2271	1
B_impute[9]	0.51	0.87	-0.91	1.88	2877	1
B_impute[10]	-0.45	0.89	-1.84	0.94	2668	1
B_impute[11]	-0.28	0.88	-1.69	1.16	2404	1
B_impute[12]	0.15	0.91	-1.33	1.60	2483	1

Compared to complete-cases

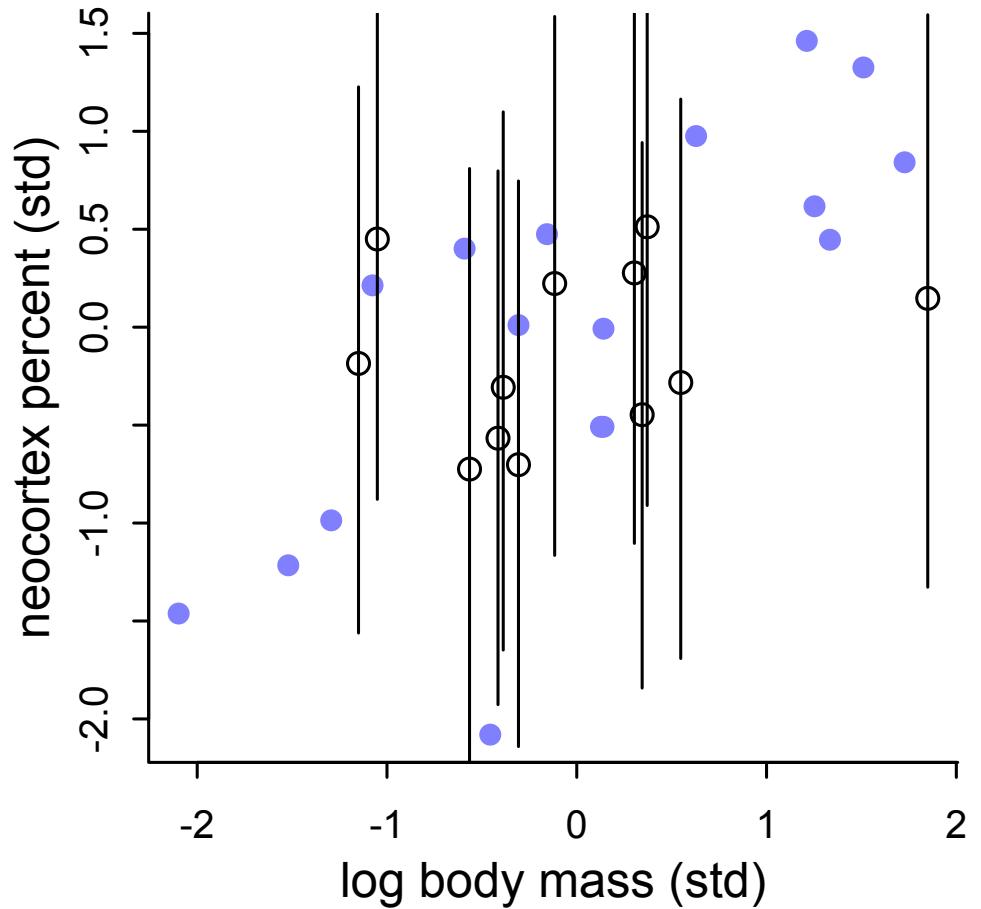
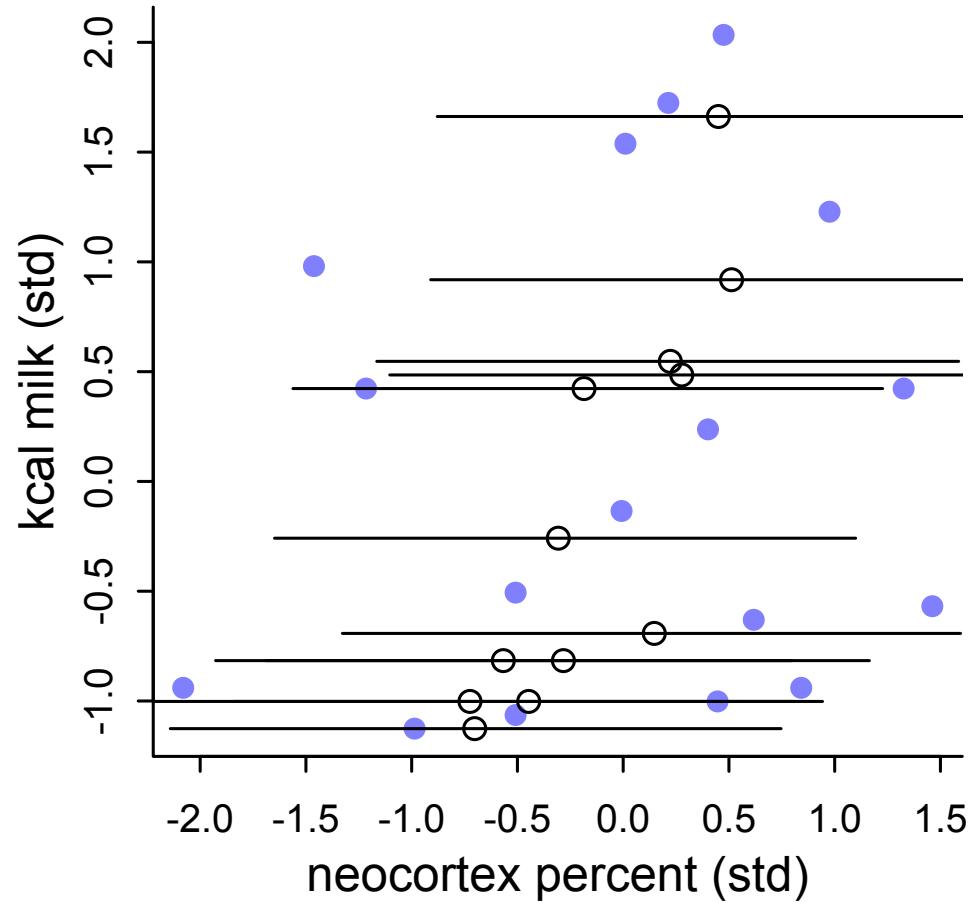


m15.3: full sample (with imputation)

m15.4: complete-cases only



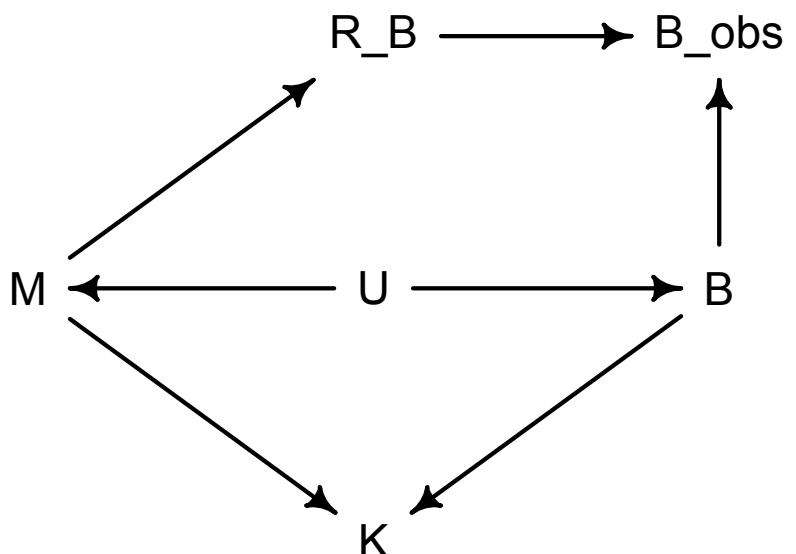
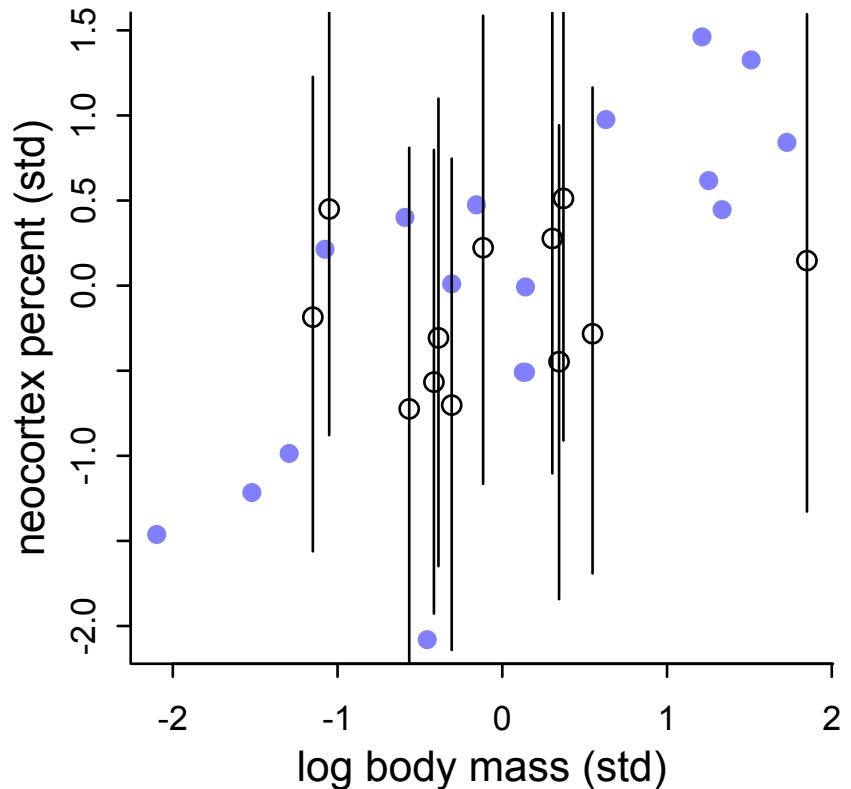
Imputed values track regression trend



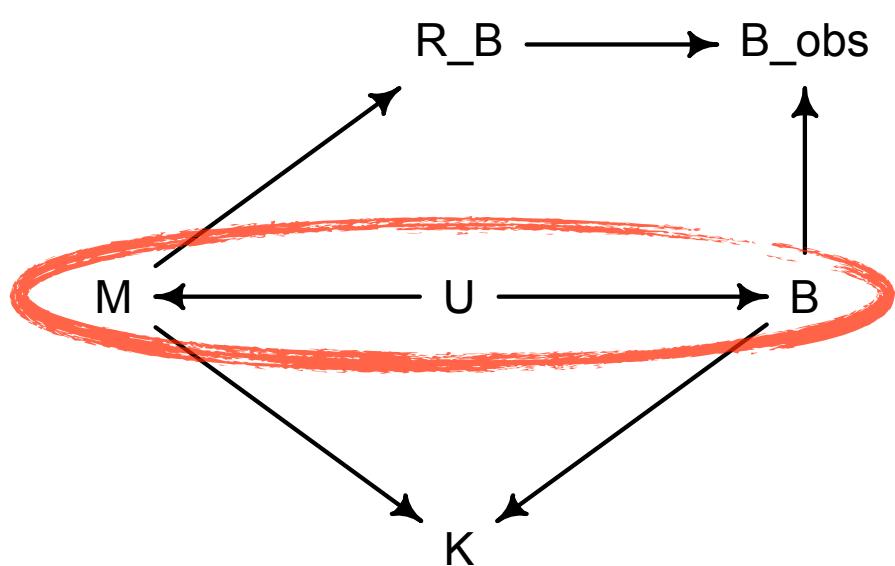
Imputed values do not track other predictor!

Results

- Observed neocortex positively associated with observed body mass
- Imputed neocortex NOT associated with observed body mass
- Can do better
 - Imputation model can use full causal model



Full Flavor Imputation



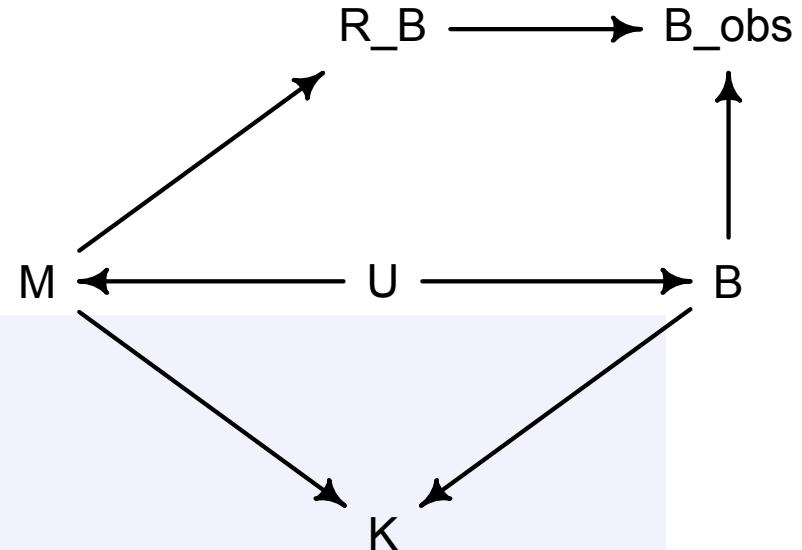
$$K_i \sim \text{Normal}(\mu_i, \sigma)$$

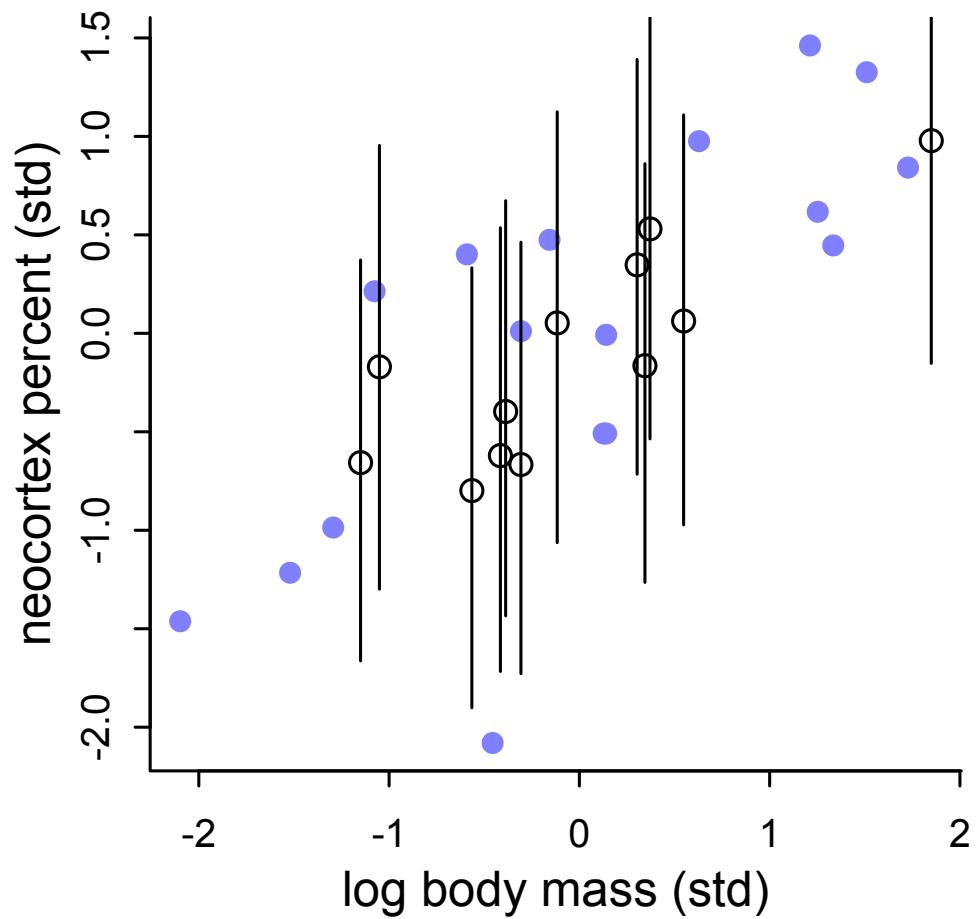
$$\mu_i = \alpha + \beta_B B_i + \beta_M \log M_i$$

$$(M_i, B_i) \sim \text{MVNormal}((\mu_M, \mu_B), \mathbf{S})$$

Full Flavor Imputation

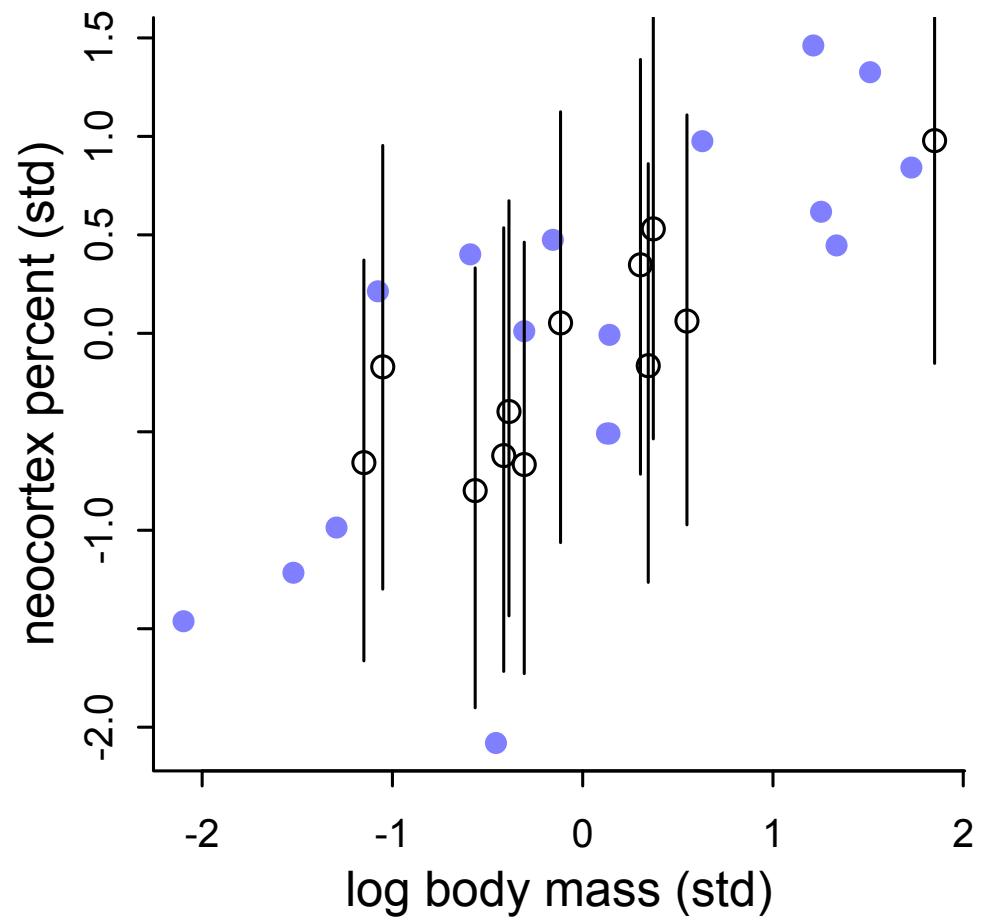
```
m15.5 <- ulam(  
  alist(  
    # K as function of B and M  
    K ~ dnorm( mu , sigma ),  
    mu <- a + bB*B_merge + bM*M,  
  
    # M and B correlation  
    MB ~ multi_normal( c(muM,muB) , Rho_BM , Sigma_BM ),  
    matrix[29,2]:MB <- append_col( M , B_merge ),  
  
    # define B_merge as mix of observed and imputed values  
    vector[29]:B_merge <- merge_missing( B , B_impute ),  
  
    # priors  
    c(a,muB,muM) ~ dnorm( 0 , 0.5 ),  
    c(bB,bM) ~ dnorm( 0, 0.5 ),  
    sigma ~ dexp( 1 ),  
    Rho_BM ~ lkj_corr(2),  
    Sigma_BM ~ exponential(1)  
  ) , data=dat_list , chains=4 , cores=4 )
```



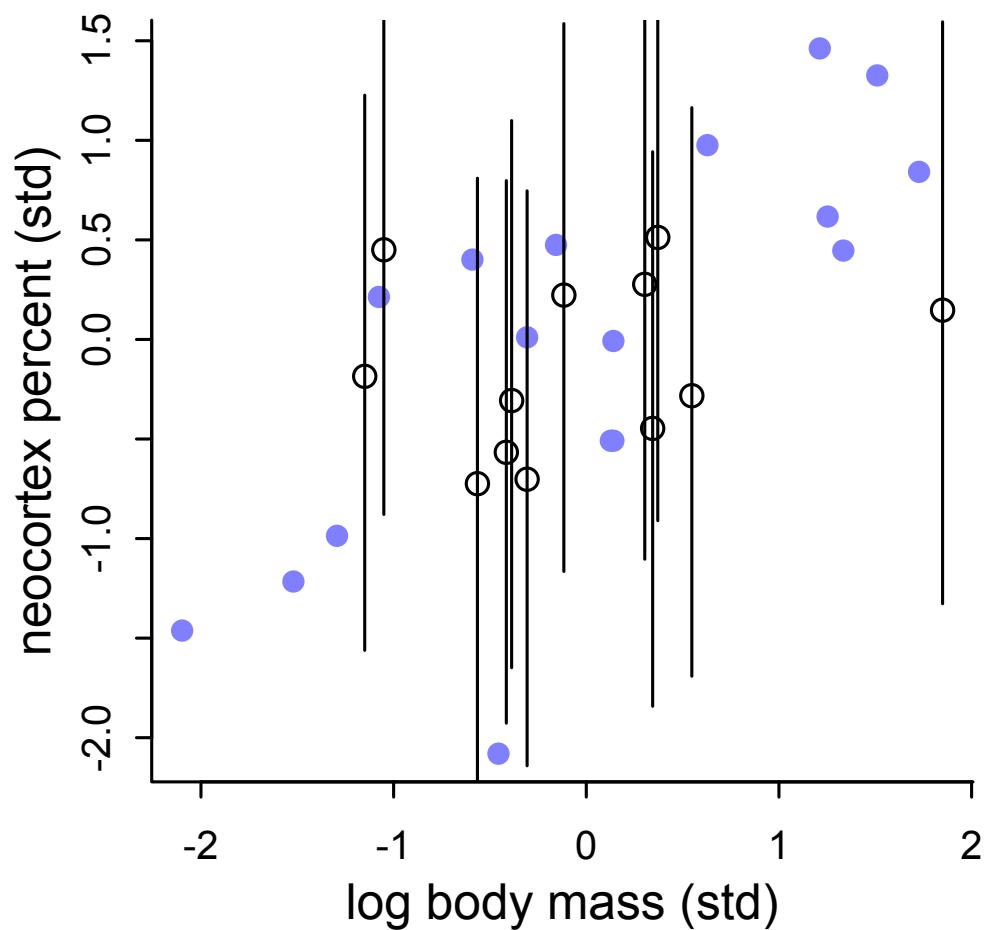


	mean	sd	5.5%	94.5%
bM	-0.65	0.22	-1.00	-0.30
bB	0.58	0.26	0.16	0.99
Rho_BM[1,1]	1.00	0.00	1.00	1.00
Rho_BM[1,2]	0.60	0.13	0.37	0.78
Rho_BM[2,1]	0.60	0.13	0.37	0.78
Rho_BM[2,2]	1.00	0.00	1.00	1.00

MVNormal

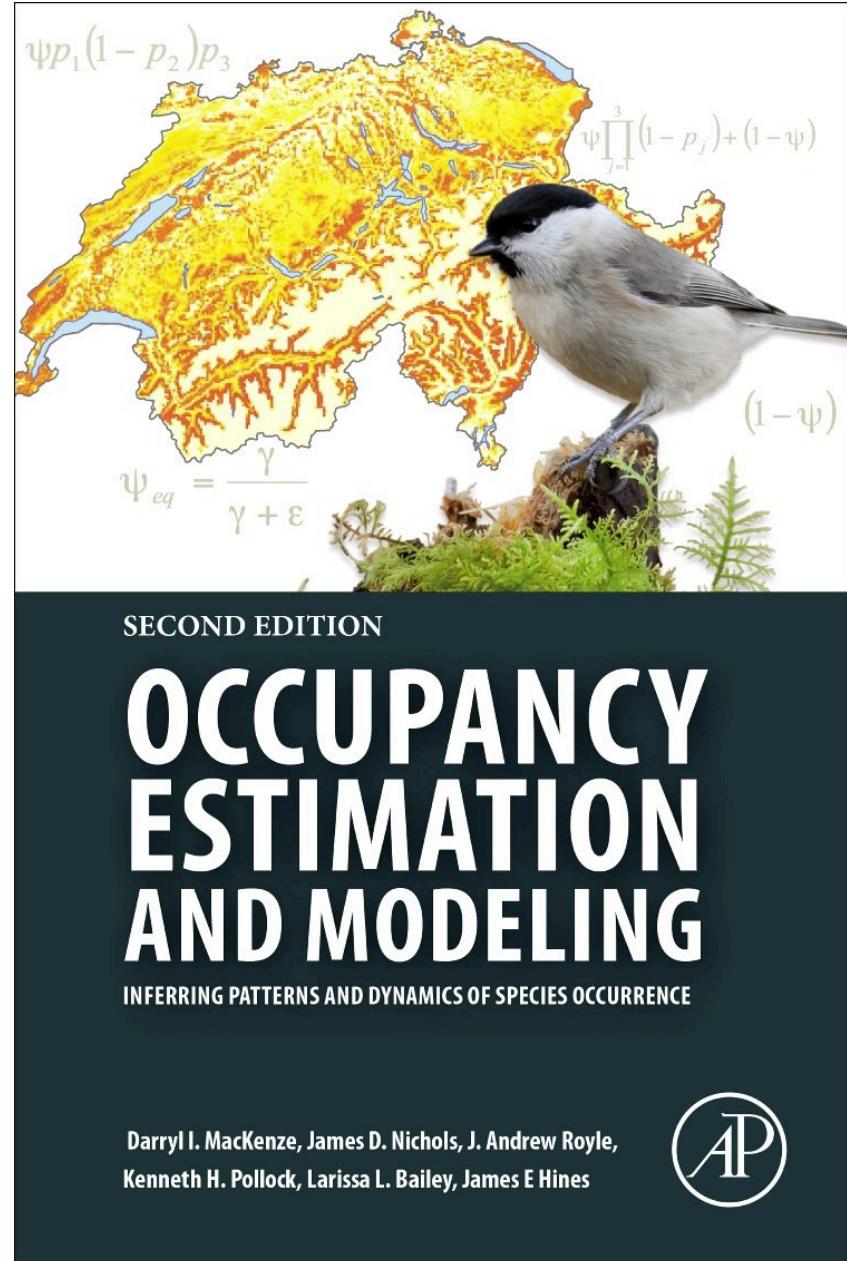


Normal



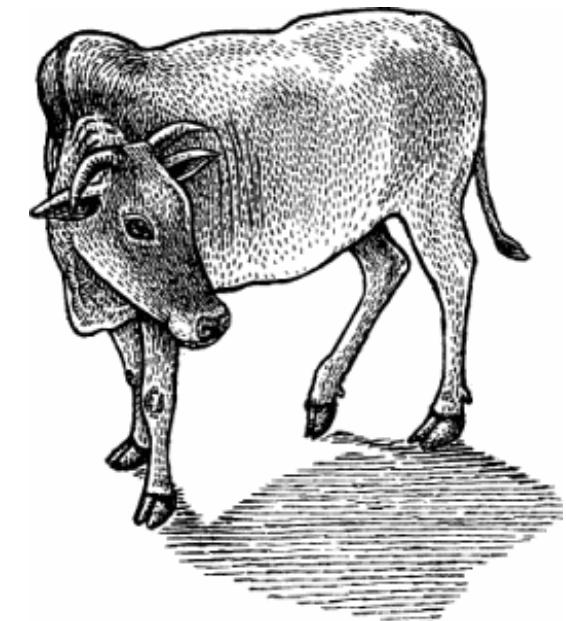
Missing data

- Can also impute discrete values, but need another technique (see text)
- Extends to many model types:
 - Mark-recapture, occupancy (presence/absence)
 - Latent-state models (hidden Markov models)



Final Homework

- A little imputation practice
- Finish for complete sense of accomplishment



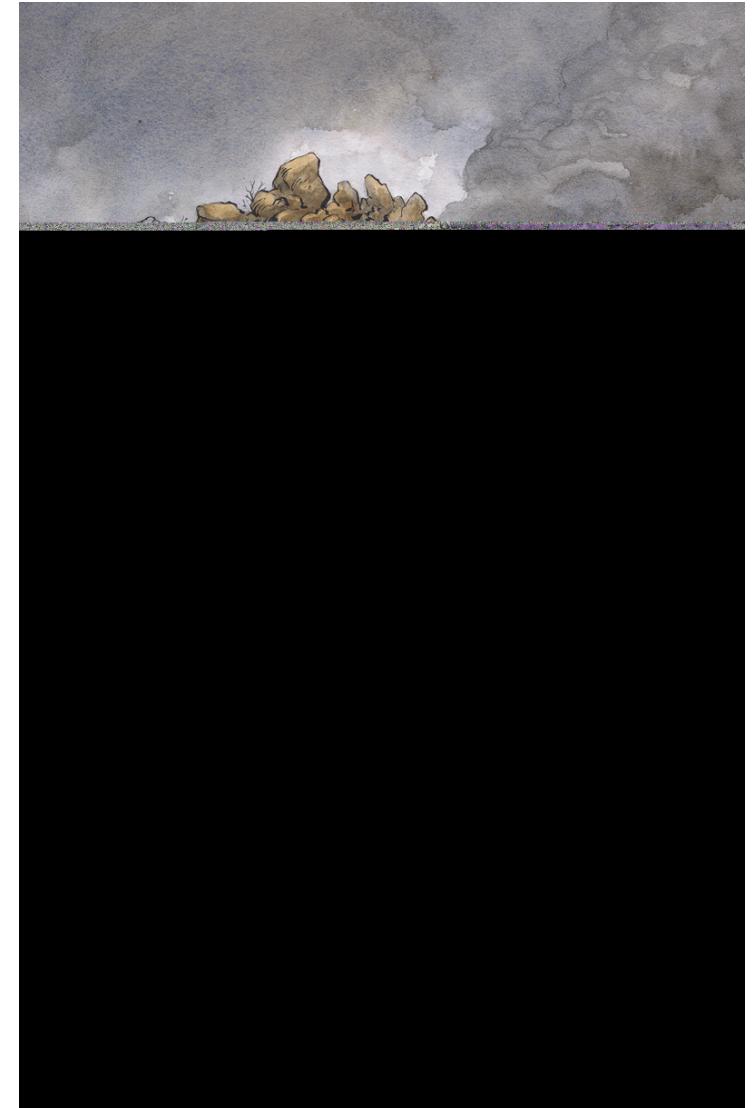
Bos primigenius

The Golem of Prague

“Even the most perfect of Golem, risen to life to protect us, can easily change into a destructive force. Therefore let us treat carefully that which is strong, just as we bow kindly and patiently to that which is weak.”



Rabbi Judah Loew ben
Bezalel (1512–1609)



From *Breath of Bones: A Tale of the Golem*